

Estimating Mutual Information from Average Classification Error

Charles Zheng and Yuval Benjamini

March 18, 2016

Abstract

Multivariate pattern analyses approaches in neuroimaging are fundamentally concerned with investigating the quantity and type of information processed by various regions of the human brain; typically, estimates of classification accuracy are used to quantify information. While a extensive and powerful library of methods can be applied to train and assess classifiers, it is not always clear how to use the resulting measures of classification performance to draw scientific conclusions: e.g. for the purpose of evaluating redundancy between brain regions. An additional confound for interpreting classification performance is the dependence of the error rate on the number and choice of distinct classes obtained for the classification task. In contrast, mutual information is a quantity which is in principle defined independently of the experimental design, and has ideal properties for comparative analyses. One obstacle to wider use of mutual information is the difficulty of estimating mutual information in high-dimensional neuroimaging data. We propose a new estimator of mutual information based on the concept of "average Bayes error." By assuming that the classes obtained in the data are drawn randomly from a continuous population, we find that the expected Bayes error converges to a function of the mutual information in a high-dimensional limit. We demonstrate the utility of our approach in simulated and real-data examples.

1 Introduction

Many of the key questions motivating neuroimaging studies involve some concept of information: which brain region processes sensory information? how is information encoded in the form of memories? To answer whether or not a particular brain region receives information about a given stimulus, it suffices to demonstrate a systematic dependence between brain activity in that region and stimulus features; this was the approach taken for studying single-cell recordings. Moving beyond single-cell recordings, it becomes desirable to use methods which can quantify the information encoded by populations of cells.

Following the seminal work by Haxby (2001), the dominant approach in this area, known as “multivariate pattern analysis” (MVPA), is to quantify information in the form of classification error. The classification approach contrasts with more traditional measures of information, such as mutual information and Fisher information. For example, to demonstrate that a particular brain region responds to a certain type of sensory information, one employs supervised learning to build a classifier which classifies the stimulus class from the brain activation in that region, and tests the statistical hypothesis that the classifier has above-chance classification accuracy. In principle, one could just as well test the statistical hypothesis that the Fisher information or mutual information between the stimulus and the activation patterns is nonzero. But in practice, the machine learning approach enjoys several advantages over approaches based on estimating Fisher information or mutual information. First, one does not need a parameterization of the stimulus space: this is highly relevant for studying complex stimuli such as faces. Secondly, the machine learning approach does not require parametric model assumptions, unlike Fisher information. Thirdly, the machine learning approach scales well with dimensionality of both the stimulus space and the response space. In contrast, nonparametric methods for estimating mutual information scale extremely poorly with dimensionality of the response space. Meanwhile, a major detractor for using classification error as a measure of information is that the classification accuracy depends on the particular choice of stimuli exemplars employed in the study and the number of partitions used to define the classes for the classification task. The difficulty of the classification task depends on the number of classes defined: high classification accuracy can be achieved relatively easily by using a coarse partition of stimuli exemplars into classes. In a meta-analysis on visual decoding, Coutanche

et al (2016) quantified the strength of a classification study using the formula

$$\text{decoding strength} = \frac{\text{accuracy} - \text{chance}}{\text{chance}}.$$

Such an approach may compensate for the differences in accuracy due purely to choice of number of classes defined; however, no theory is provided to justify the formula.

In contrast, mutual information has ideal properties for the quantitatively comparing information between different studies, or between different brain regions, subjects, featurization models, or modalities. Not only is the mutual information defined independently of the arbitrary definition of stimulus classes (albeit still dependent on an implied distribution over stimuli), it is even meaningful to discuss the difference between the mutual information measured for one system and the mutual information for a second system. Due to these appealing properties, as well as the success of information theoretic approaches in more traditional neuroscience studies, mutual information has been frequently proposed for application in neuroimaging studies (Quiroga 2009). However, since nonparametric approaches for estimating mutual information are unusably inaccurate for typical high-dimensional data sets, alternative approaches must be employed. One can tractably estimate mutual information by assuming a multivariate Gaussian model: however, this approach essentially assumes a linear relationship between the input and output, and hence fails to quantify nonlinear dependencies. Alternatively, one can obtain a lower bound on the mutual information from the confusion matrix of the classifier. This is the most popular approach for estimating mutual information in neuroimaging studies, but suffers from known shortcomings (Gastpar 2010, Quiroga 2009).

The idea of linking classification performance to mutual information is not new: after all, Shannon’s original motivation was to characterize the minimum achievable error probability of a noisy communication channel. More explicitly, Fano’s inequality provides a lower bound on mutual information in relation to the optimal prediction error, or Bayes error. Fano’s inequality can be further refined to obtain a tighter lower bound on mutual information (Tebbe and Dwyer 1968.)

In practice, the idea of deriving estimates of mutual information from classification performance occupies an advantageous middle ground between the two extremes of nonparametric and parametric approaches for estimating mutual information. In neuroimaging data, we lack prior knowledge for spec-

ifying parametric models, and the data is too high-dimensional for nonparametric approaches, but we have a sufficient idea of the general “structure” in the data to achieve above-chance classification rates.

However, there are several practical issues for estimating mutual information from classification accuracy. One, the achievable classification accuracy depends on the amount of data available for training. The test error is an unbiased estimate of the true error rate of the classifier on new data (called generalization error); but in turn, the generalization error can only be meaningfully interpreted as a lower bound on the error rate of the ideal classifier, or Bayes error. Secondly, even if one could obtain the Bayes error, the Bayes error itself depends on the choices made by the experimenter in regards to the stimuli exemplar chosen in the experiment, and the decision of how to partition those exemplars in the classification task. For example, Nishimoto et al classified segments of a movie clips based on activation patterns, but the definition of the classification task, and the achievable classification performance, depends not only on the particular movie clips used in the experiment, but also the choice of time interval used to define discrete classes: defining each class to be a 1sec segment of movie results in more distinct classes and lower classification accuracy than defining each class to be a 4sec segment of movie. The Bayes error, and any estimate of mutual information based on the Bayes error, would therefore be necessarily dependent on the experimental parameters.

In principle, the first issue can be overcome by having sufficiently many observations: but a practical issue is that an astronomical amount of data might be needed to learn the optimal classification rule. A more efficient approach would be to estimate the Bayes error from the classification performance. Since the Bayes error is the large-sample limit of the achieved classification error, a promising approach is to perform classification using differently sized subsamples of the training data, producing a plot of classification error versus sample size—a “learning curve.” One can then extrapolate the learning curve to estimate the Bayes error (Cortes et al. 1994.) However, much work remains to develop rigorous methodology for estimating Bayes error, and so we leave this first issue for future work.

The starting point for our methodology is a proposal for overcoming the second issue: the non-uniqueness of the Bayes error. We define a notion of *K-class average Bayes error* which is uniquely defined for any given stimulus distribution and stochastic mapping from stimulus to response. The *K-class average Bayes error* is simply the expectation of the Bayes error when

K stimuli exemplars are drawn i.i.d. from the stimulus distribution, and treated as distinct classes. Hence the average Bayes error can in principle be estimated if the appropriate randomization is employed for designing the experiment.

While the K -class average Bayes error is defined independently of the particular choice of stimuli, the quantity still depends on the choice of number of classes, K . In comparison, the mutual information provides a quantification of information which does not depend on K , allowing more flexible comparisons and easier interpretation. Hence our main theoretical contribution is the derivation of an asymptotic relationship between the K -class average Bayes error and the mutual information, which in practice means that any estimate of the K -class Bayes error can be converted into an estimate of mutual information; and the resulting estimator of mutual information should be asymptotically independent of the choice of number of classes K .

2 Average Bayes Error

The following simplified model captures the essence of many neuroimaging studies. Let \mathcal{X} be a space of stimuli, represented by q -dimensional vectors. In the design stage of the experiment, a set of stimuli exemplars $\{x_1, \dots, x_k\}$ are selected, and then one specifies a T -length sequence of stimuli $(x_{i_1}, \dots, x_{i_n})$ to be presented to the subject. In the execution of the experiment, an activation pattern, or *response* y_t is obtained for each of the stimuli presentations x_{i_t} . Generally y_t is a p -dimensional vector, representing activity levels of p disjoint brain regions. To simplify, assume that each of the k stimuli is presented a total of r times in the sequence; further assume that the responses to each stimulus presentation are conditionally independent, hence the ordering of the sequence does not matter. Henceforth we can let y_i^j denote the response to the j th presentation of the stimulus x_i .

We now restate the model in application-independent terms. Let $\mathcal{X} \subset \mathbb{R}^q$ and $\mathcal{Y} \subset \mathbb{R}^p$; let $p(x)$ be a probability density on \mathcal{X} . For every $x \in \mathcal{X}$, let $p_x(y)$ be a probability density on \mathcal{Y} . Define the joint density

$$p(x, y) = p(x)p_x(y)$$

so that $p_x(y)$ can also be written $p(y|x)$. Take a subset $\{x_1, \dots, x_k\} \subset \mathcal{X}$. Let Y_i^j be a random vector distributed according to density $p_{x_i}(y)$; and let Y_i^j be conditionally independent given $\{x_1, \dots, x_k\}$ for $i = 1, \dots, k$ and $j = 1, \dots, r$.

In MVPA, one carries out a classification task to assess whether y contains information about x . Formally, a classification rule is any (possibly stochastic) mapping $f : \mathcal{Y} \rightarrow \{1, \dots, k\}$. The *generalization error* of the classification rule is

$$e_{gen}(f) = \frac{1}{k} \sum_{i=1}^k \Pr[f(Y) \neq i | X = x_i].$$

A trivial classification rule which outputs the result of a k -sided die roll for all inputs y would achieve a generalization error of $e_{gen} = \frac{1}{k}$. Conversely, even a single counterexample with $e_{gen} < \frac{1}{k}$ is indicative that y contains nonzero information about x . Hence, in order to demonstrate that y is informative of x , one tests the null hypothesis

$$H_0 : e_{gen}(f) = \frac{1}{k}.$$

Rejecting the null hypothesis for a given classification rule f can be taken as evidence that y is informative of x . In the discussion thus far, “information” refers to an informal, intuitive notion, but it is also possible to formally establish that violation of the null hypothesis implies nonzero mutual information $I(X; Y)$.

The classifier is a functional which maps a set of observations to a classification rule,

$$\mathcal{F} : \{(x_1, y_1), \dots, (x_m, y_m)\} \mapsto f(\cdot);$$

informally, a classifier is an algorithm which “learns” a classification rule from data. For a valid test of H_0 , either the data-splitting approach or the cross-validation approach can be used. In the data-splitting approach, one creates a *training set* consisting of r_1 repeats per class,

$$\{(x_1, y_1^1), \dots, (x_1, y_1^{r_1}), \dots, (x_m, y_m^1), \dots, (x_m, y_m^{r_1})\}$$

and a *test set* consisting of the remaining $r_2 = r - r_1$ repeats.

$$\{(x_1, y_1^{r_1+1}), \dots, (x_1, y_1^r), \dots, (x_m, y_m^{r_1+1}), \dots, (x_m, y_m^r)\}.$$

In the cross-validation approach, one applies the data-splitting approach multiple times, with different training and test partitions, and aggregates the results. We further discuss the cross-validation approach in section X.

In the data-splitting approach, one obtains the classification rule f by applying the classifier to the training data,

$$f = \mathcal{F}(\{(x_1, y_1^1), \dots, (x_1, y_1^{r_1}), \dots, (x_m, y_m^1), \dots, (x_m, y_m^{r_m})\})$$

The test statistic of interest is the test error, defined as

$$e_{test} = \frac{1}{kr_2} \sum_{i=1}^k \sum_{j=r_1+1}^r \mathbf{I}(f(y_i^j) \neq i).$$

Due to the conditional independence of the training set and test set, e_{test} is an unbiased estimate of e_{gen} . Hence various approaches can be used to obtain a threshold c_α such that $\Pr[e_{test} < c_\alpha] \leq \alpha$ holds (approximately) under the null hypothesis. To name a few, one can employ permutation tests, tests based on a universal variance bound, or the generalized likelihood ratio test. Regardless of the procedure used to derive the threshold c_α , The hypothesis H_0 is rejected at level α if $e_{test} < c_\alpha$.

While tests of the generalization error suffice to establish the presence of information, the generalization error is less satisfactory as a measure of the information between X and Y , since the e_{gen} depends on the classification rule f obtained, and hence, varies randomly depending on the sampling of the data. For most reasonable classifiers, the generalization error decreases as the number of training observations increases. Therefore, the sample size is an additional confounding factor for interpreting the generalization error.

A more ideal measure of information, still related to the classification, is the Bayes error, which is simply the *optimal* generalization error

$$e_{Bayes} = \min_f e_{gen}(f).$$

Due to Bayes' theorem, the optimal classification rule f^* which achieves the Bayes error can be given explicitly: it is the maximum a posteriori (MAP) rule

$$f^*(y) = \operatorname{argmax}_{i=1}^k p(y|x_i).$$

Estimating e_{Bayes} is much more difficult in practice than estimating e_{gen} ; however, nonparametric estimators of e_{Bayes} have been proposed in the literature. Cover (1969) first noted that the leave-one-out error of 1-nearest neighbor, e_{1nn} , asymptotically bounds the Bayes error in the sense that

$$\frac{1}{2} \mathbf{E}[e_{1nn}] \leq e_{Bayes} \leq \mathbf{E}[e_{1nn}].$$

Later, Fukunaga and Kessel (1971) and Fralick and Scott (1971) both proposed a consistent estimator of e_{Bayes} using the test error from kernel density estimation-based classifiers. Fukunaga and Hostetler 1975 improved on the density estimation approach by estimating the generalization error directly, rather than using empirical test error (hence reducing variance.)

Nevertheless, the convergence rates of nonparametric methods of Bayes risk estimation are too slow to provide much assurance for their use in high-dimensional problems. On the other hand, the empirical success of supervised learning approaches seems to justify a certain optimism that at least one of the popular black-box methods (random forests, neural networks, kernel SVM) can come close to the Bayes error rate within realistic sample sizes. Methods such as random forests and neural networks are known to have *universal approximation* properties that guarantee consistent recovery of the classification rule, given that the model complexity is scaled at an appropriate rate as the sample size increases; but methods such as kernel SVM have much more restrictive function classes and therefore have no guarantee of being able to recover the Bayes error. Yet, it is seen in many learning tasks that the test error is very similar for very different classification methods, some universal and some non-universal, suggesting that the true signal lies in a low-complexity function class which can be adequately approximated using a variety of methods. Empirical evidence for these phenomenon are collected in studies on the learning curves of classifiers, beginning with Cortes 1994, and applied to specific applications by Figueroa et al. 2012, Beleites et al. 2013. We leave further discussion of Bayes error estimation for future work; in the theoretical portion of the current paper, we will simply leave unspecified the method used to estimate Bayes error, and in our simulation and real data examples, we will use an ad hoc approach for estimating Bayes error from test error for the purpose of demonstrating our method.

As we pointed out in the introduction, the quantity e_{Bayes} is not a parameter of the joint distribution $p(x, y)$, but rather depends on the specific exemplars x_1, \dots, x_k selected; hence, one may write $e_{Bayes}(x_1, \dots, x_k)$ to emphasize this dependence. In many studies, the end goal is to make some inference about the joint distribution $p(x, y)$, such as the mutual information $I(X; Y)$. However, an estimate of e_{Bayes} cannot possibly be sufficient as a basis for inference about $p(x, y)$ because there is no direct link between e_{Bayes} and $p(x)$, given that we made no assumption on how x_1, \dots, x_k were

obtained. This motivates definition of the *average* Bayes error,

$$e_{ABE} = \mathbf{E}[e_{Bayes}(X_1, \dots, X_k)],$$

where X_1, \dots, X_K are drawn i.i.d. from $p(x)$. The k -class average Bayes error is uniquely defined for any joint distribution $p(x, y)$, and supposing that an unbiased estimator for Bayes error exists, \hat{e}_{Bayes} , then \hat{e}_{Bayes} will also be unbiased for e_{ABE} supposing that the experimental design randomized the selection of X_1, \dots, X_k by drawing the vectors independently from $p(x)$.

In the following section we investigate the relationship between e_{ABE} and the mutual information $I(X; Y)$.

3 Asymptotic Theory

We give the proof with minimal context or motivation. See the paper for more details.

Fix integer $K \geq 2$. Let $p^{[d]}(x, y)$ be a sequence of probability density functions, where x is of dimension $p^{[d]}$ and y is of dimension $q^{[d]}$. Let $p^{[d]}(x)$ and $p^{[d]}(y)$ denote the marginal densities, and let

$$p^{[d]}(y|x) = p^{[d]}(x, y)/p^{[d]}(x).$$

Let $(X^{([d], i)}, Y^{([d], i)})$ be iid random variates from $p^{[d]}(x, y)$ for $i = 0, \dots, K-1$; we will suppress the superscripts $[d]$ and/or (i) when convenient. Recall the definitions of entropy,

$$H(X) = - \int p(x) \log p(x) dx,$$

and mutual information

$$I(X; Y) = \int p(x, y) \log \frac{p(x, y)}{p(x)p(y)} dx dy.$$

Furthermore, define the K -class average Bayes error as

$$ABE_K \Pr[p(Y^{(0)}|X^{(0)}) < \max_{i=1}^{K-1} p(Y^{(0)}|X^{(i)})].$$

Define

$$u^{[d]}(x, y) = \log p^{[d]}(x, y) - \log p^{[d]}(x) - \log p^{[d]}(y),$$

and define

$$\ell_{ij}^{[d]} = \log p(y^{(i)}|x^{(j)})$$

Theorem. Let $p^{[d]}(x, y)$ be a sequence of joint densities for $d = 1, 2, \dots$ as given above. Further assume that

A1. $\lim_{d \rightarrow \infty} I(X^{[d]}; Y^{[d]}) = \iota < \infty$.

A2. There exists a sequence of scaling constants $a_{ij}^{[d]}$ and $b_{ij}^{[d]}$ such that the random vector $(a_{ij}\ell_{ij}^{[d]} + b_{ij}^{[d]})_{i,j=1,\dots,K}$ converges in distribution to a multivariate normal distribution.

A3. There exists a sequence of scaling constants $a^{[d]}, b^{[d]}$ such that

$$a^{[d]}u(X^{(1)}, Y^{(2)}) + b^{[d]}$$

converges in distribution to a univariate normal distribution.

A4. For all $i \neq k$,

$$\lim_{d \rightarrow \infty} \text{Cov}[u(X^{(i)}, Y^{(j)}), u(X^{(k)}, Y^{(j)})] = 0.$$

Then for ABE_K as defined above, we have

$$\lim_{d \rightarrow \infty} \text{ABE}_K = \pi_K(\iota)$$

where

$$\pi_K(\iota) = 1 - \int_{\mathbb{R}} \phi(z - \sqrt{2\iota})(1 - \Phi(z))^{K-1} dz$$

where ϕ and Φ are the standard normal density function and cumulative distribution function, respectively.

Proof.

For $i = 1, \dots, K-1$, define

$$Z_i = \log p(Y^{(0)}|X^{(i)}) - \log p(Y^{(0)}|X^{(0)}).$$

Then, we claim that $\vec{Z} = (Z_1, \dots, Z_{K-1})$ converges in distribution to

$$\vec{Z} \sim N \left(-2\iota, \begin{bmatrix} 4\iota & 2\iota & \cdots & 2\iota \\ 2\iota & 4\iota & \cdots & 2\iota \\ \vdots & \vdots & \ddots & \vdots \\ 2\iota & 2\iota & \cdots & 4\iota \end{bmatrix} \right).$$

Combining the claim with the lemma in the following section yields the desired result.

To prove the claim, it suffices to derive the limiting moments

$$\mathbf{E}[Z_i] \rightarrow -2\iota,$$

$$\text{Cov}[Z_i] \rightarrow 4\iota,$$

$$\text{Cov}[Z_i, Z_j] \rightarrow 2\iota,$$

for $i \neq j$, since then assumption A2 implies the existence of a multivariate normal limiting distribution with the given moments.

Before deriving the limiting moments, note the following identities. Let $X' = X^{(1)}$ and $Y = Y^{(0)}$.

$$\mathbf{E}[e^{u(X', Y)}] = \int p(x)p(y)e^{u(x,y)}dxdy = \int p(x,y)dxdy = 1.$$

Therefore, from assumption A3 and the formula for gaussian exponential moments, we have

$$\lim_{d \rightarrow \infty} \mathbf{E}[u(X', Y)] - \frac{1}{2} \text{Var}[u(X', Y)] = 0.$$

Let $\sigma^2 = \lim_{d \rightarrow \infty} \text{Var}[u(X', Y)]$. Meanwhile, by applying assumption A2,

$$\begin{aligned} \lim_{d \rightarrow \infty} I(X; Y) &= \lim_{d \rightarrow \infty} \int p(x, y)u(x, y)dxdy = \lim_{d \rightarrow \infty} \int p(x)p(y)e^{u(x,y)}u(x, y)dxdy \\ &= \lim_{d \rightarrow \infty} \mathbf{E}[e^{u(X, Y')}u(X, Y')] \\ &= \int_{\mathbb{R}} e^z z \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(z+\sigma^2/2)^2}{2\sigma^2}} dz \quad (\text{applying A2}) \\ &= \int_{\mathbb{R}} z \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(z-\sigma^2/2)^2}{2\sigma^2}} dz \\ &= \frac{1}{2}\sigma^2. \end{aligned}$$

Therefore,

$$\sigma^2 = 2\iota,$$

and

$$\lim_{d \rightarrow \infty} \mathbf{E}[u(X', Y)] = -\iota.$$

Once again by applying A2, we get

$$\begin{aligned}
\lim_{d \rightarrow \infty} \text{Var}[u(X, Y)] &= \lim_{d \rightarrow \infty} \int (u(x, y) - \iota)^2 p(x, y) dx dy \\
&= \lim_{d \rightarrow \infty} \int (u(x, y) - \iota)^2 e^{u(x, y)} p(x) p(y) dx dy \\
&= \lim_{d \rightarrow \infty} \mathbf{E}[(u(X', Y) - \iota)^2 e^{u(X', Y)}] \\
&= \int (z - \iota)^2 e^z \frac{1}{\sqrt{4\pi\iota}} e^{-\frac{(z+\iota)^2}{4\iota}} dz \quad (\text{applying A2}) \\
&= \int (z - \iota)^2 \frac{1}{\sqrt{4\pi\iota}} e^{-\frac{(z-\iota)^2}{4\iota}} dz \\
&= 2\iota.
\end{aligned}$$

We now proceed to derive the limiting moments. We have

$$\begin{aligned}
\lim_{d \rightarrow \infty} \mathbf{E}[Z] &= \lim_{d \rightarrow \infty} \mathbf{E}[\log p(Y|X') - \log p(Y|X)] \\
&= \lim_{d \rightarrow \infty} \mathbf{E}[u(X', Y) - u(X, Y)] = -2\iota.
\end{aligned}$$

Also,

$$\begin{aligned}
\lim_{d \rightarrow \infty} \text{Var}[Z] &= \lim_{d \rightarrow \infty} \text{Var}[u(X', Y) - u(X, Y)] \\
&= \lim_{d \rightarrow \infty} \text{Var}[u(X', Y)] + \text{Var}[u(X, Y)] \quad (\text{using assumption A4}) \\
&= 4\iota,
\end{aligned}$$

and similarly

$$\begin{aligned}
\lim_{d \rightarrow \infty} \text{Cov}[Z_i, Z_j] &= \lim_{d \rightarrow \infty} \text{Var}[u(X, Y)] \quad (\text{using assumption A4}) \\
&= 2\iota.
\end{aligned}$$

This concludes the proof. \square .

3.1 Lemma

Lemma. Suppose $(Z_*, Z_1, \dots, Z_{K-1})$ are jointly multivariate normal, with $\mathbf{E}[Z_* - Z_1] = \alpha$, $\text{Var}(Z_*) = \beta$, $\text{Cov}(Z_*, Z_i) = \gamma$, $\text{Var}(Z_i) = \delta$, and $\text{Cov}(Z_i, Z_j) = \epsilon$ for all $i, j = 1, \dots, K-1$. Then, letting

$$\mu = \frac{\mathbf{E}[Z_* - Z_i]}{\sqrt{\frac{1}{2} \text{Var}(Z_i - Z_j)}} = \frac{\alpha}{\sqrt{\delta - \epsilon}},$$

$$\nu^2 = \frac{\text{Cov}(Z_* - Z_i, Z_* - Z_j)}{\frac{1}{2} \text{Var}(Z_i - Z_j)} = \frac{\beta + \epsilon - 2\gamma}{\delta - \epsilon},$$

we have

$$\begin{aligned} \Pr[Z_* < \max_{i=1}^{K-1}] &= \Pr[W < M_{K-1}] \\ &= 1 - \int -\frac{1}{\sqrt{2\pi\nu^2}} e^{-\frac{(w-\mu)^2}{2\nu^2}} (1 - \Phi(w))^{K-1} dw, \end{aligned}$$

where $W \sim N(\mu, \nu^2)$ and M_{K-1} is the maximum of $K-1$ independent standard normal variates, which are independent of W .

[ADD PROOF OF LEMMA HERE]