Information Theory Notes

Charles Zheng and Yuval Benjamini

December 5, 2015

These are preliminary notes.

1 Classification in high-dimension, fixed SNR regime

We observe a data point y_* which belongs to one of K classes. The distribution in the ith class is $N(\mu_i, \Omega)$. We have another dataset with r repeats per class, which we use to estimate the centroids μ_i: we obtain estimates $\hat{\mu}_i \sim N(\mu_i, r^{-1}\Omega)$. The class centroids were originally drawn i.i.d. from a multivariate normal $N(0, I)$. Furthermore Ω is unknown and have to be estimated as well: assume we have obtained estimate $\hat{\Omega}$ via some method. Without loss of generality, take the Kth class to be the true class of y_*. Write $\hat{\mu}_* = \hat{\mu}_K$.

The classification rule is given by

$$\text{Estimated class} = \text{argmin}_i \ (y_* - B\hat{\mu}_i)^T A (y_* - B\hat{\mu}_i)$$

where A and B are matrices based on $\hat{\Omega}$. The Bayes rule is given by

$$A_{Bayes} = (I + \Omega - (I + r^{-1}\Omega)^{-1})^{-1}$$

$$B_{Bayes} = (I + r^{-1}\Omega)^{-1}.$$

The "plug-in" estimates of A and B are

$$A = (I + \hat{\Omega} + (I + r^{-1}\hat{\Omega})^{-1})^{-1}$$

$$B = (I + r^{-1}\hat{\Omega})^{-1}.$$

Note that

$$(y_* - B\hat{\mu}_i)^T A(y^* - B\hat{\mu}_i) = ||A^{1/2}y_* - A^{1/2}B\hat{\mu}_i||^2.$$

Therefore the classification rule is

$$\text{Estimated class} = \text{argmin}_i \; Z_i,$$

where

$$Z_i = ||A^{1/2}y_* - A^{1/2}B\hat{\mu}_i||^2.$$

We have

$$\begin{bmatrix} A^{1/2}y \\ A^{1/2}B\hat{\mu}_* \\ A^{1/2}B\hat{\mu}_i \end{bmatrix} \sim N\left(\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} A^{1/2}(I+\Omega)A^{1/2} & A^{1/2}BA^{1/2} & 0 \\ & A^{1/2}B(I+\frac{\Omega}{r})BA^{1/2} & 0 \\ & & A^{1/2}B(I+\frac{\Omega}{r})BA^{1/2} \end{bmatrix} \right)$$

Therefore

$$\mathbf{E}Z_i = \begin{cases} \text{tr}[A(I+\Omega+(B(I+r^{-1}\Omega)B))] & \text{for } i \neq K \\ \text{tr}[A(I+\Omega+(B(I+r^{-1}\Omega)B)-2B)] & \text{for } i = K \end{cases},$$

$$\text{Cov}(Z_i, Z_j) = \begin{cases} 2\text{tr}[A(I+\Omega)]^2 & \text{for } i \neq j \neq K \\ 2\text{tr}[A(I+\Omega-B)]^2 & \text{for } i = K, j \neq K \\ 2\text{tr}[A(I+\Omega+B(I+r^{-1}\Omega)B)]^2 & \text{for } i = j \neq K \\ 2\text{tr}[A(I+\Omega+B(I+r^{-1}\Omega)B-2B)]^2 & \text{for } i = j = K \end{cases}.$$

1.1 Equivalence of limiting MI and multivariate SNR

The mutual information between μ and y is given by

$$\text{MI} = -\log\det(I-(I+\Omega)^{-1}).$$

We will show that if $\text{tr}[\Omega^{-1}] \to c$, and $\text{tr}[\Omega^{-2}] \to 0$, then

$$\begin{aligned} \lim_{p\to\infty} \text{MI} &= \lim_{p\to\infty} -\log\det(I-(I+\Omega)^{-1}) \\ &= \lim_{p\to\infty} -\log\det(I-\Omega^{-1/2}(\Omega^{-1}+I)^{-1}\Omega^{-1/2}) \\ &= \lim_{p\to\infty} -\log\det(I-\Omega^{-1}+\Omega^{-2}) \\ &= \lim_{p\to\infty} -\log\det(I-\Omega^{-1}) \\ &= \lim_{p\to\infty} \log\det(I+\Omega^{-1}) \\ &= \lim_{p\to\infty} \text{tr}[\Omega^{-1}] = c \end{aligned}$$

1.2 Ω known, fixed r

Suppose

$$\hat{\Omega} = \Omega.$$

Then,

$$\mathbf{E}Z_i = \begin{cases} p + 2\mathrm{tr}[AB] & \text{for } i \neq K \\ p & \text{for } i = K \end{cases},$$

and

$$\mathrm{Cov}(Z_i, Z_j) = \begin{cases} 2\mathrm{tr}[I + AB]^2 & \text{for } i \neq j \neq K \\ 2p & \text{for } i = K, j \neq K \\ \mathrm{tr}[I + 2AB]^2 & \text{for } i = j \neq K \\ 2p & \text{for } i = j = K \end{cases}.$$

Let $\gamma_i(r)$ denote the eigenvalues of AB: we have

$$\gamma_i(r) = \frac{1}{(\omega_i^2 + \omega_i)/r + \omega_i}$$

where ω_i are the eigenvalues of Ω.

The misclassification probability is

$$\mathrm{MC} = \Pr[N(\mu(r), \sigma^2(r)) < M_{K-1}]$$

where M_{K-1} is the maximum of $K-1$ independent standard normal variates, and

$$\mu(r) = \frac{2\sum_{i=1}^p \gamma_i}{\sqrt{\sum_{i=1}^p 6\gamma_i^2 + 4\gamma_i}},$$
$$\sigma^2(r) = \frac{\sum_{i=1}^p \gamma_i^2 + 2\gamma_i}{\sum_{i=1}^p 3\gamma_i^2 + 2\gamma_i}.$$

Under the condition that

$$\mathrm{tr}[\Omega^{-1}] \to \text{const.}, \ \mathrm{tr}[\Omega^{-2}] \to 0$$

we have

$$\mu(r) \to \sqrt{\sum_{i=1}^p \gamma_i(r)},$$
$$\sigma^2(r) \to 1.$$

1.3 Known Ω, growing r

Assume that Ω/p has a limiting spectrum, i.e. defining

$$H^{(p)} = \sum_{i=1}^{p} \frac{1}{p} \delta_{\omega_i/p}$$

we have

$$H^{(p)} \to H$$

for some probability measure H on $\mathbb{R}^+$. Note then that

$$\mathrm{MI} = \lim_{p\to\infty} \mathrm{tr}[\Omega^{-1}] = \int \nu^{-1} dH(\nu).$$

Also, suppose that $\frac{r}{p} \to \kappa$. Then, define

$$\begin{aligned}
\mu^*(\kappa) = \lim_{p\to\infty} \mu(r) &= \lim_{p\to\infty} \frac{2\sum_{i=1}^{p} \gamma_i}{\sqrt{\sum_{i=1}^{p} 6\gamma_i^2 + 4\gamma_i}}, \\
&= \lim_{p\to\infty} \frac{2p \int \frac{1}{\frac{p^2\nu^2+p\nu}{r} + p\nu} dH(\nu)}{\sqrt{p\int 6\left(\frac{1}{\frac{p^2\nu^2+p\nu}{r}+p\nu}\right)^2 + 4\left(\frac{1}{\frac{p^2\nu^2+p\nu}{r}+p\nu}\right) dH(\nu)}} \\
&= \lim_{p\to\infty} \sqrt{p\int \frac{1}{\frac{p^2\nu^2+p\nu}{r} + p\nu} dH(\nu)} \\
&= \lim_{p\to\infty} \sqrt{\int \frac{1}{\frac{p\nu^2}{r} + (1+1/r)\nu} dH(\nu)} \\
&= \sqrt{\int \lim_{p\to\infty} \frac{1}{\frac{p\nu^2}{r} + (1+1/r)\nu} dH(\nu)} \\
&= \sqrt{\int \frac{1}{\frac{\nu^2}{\kappa} + \nu} dH(\nu)},
\end{aligned}$$

while for similar reasons

$$\lim_{p\to\infty} \sigma^2(r) = 1.$$

Thus, we have $\mu^*(\kappa)$ increasing in κ, and tending to

$$\mu^*(\infty) = \int \nu^{-1} H(d\nu)$$

as $\kappa \to \infty$. But, note that $\mu^*(\infty)$ is also the limiting mutual information.

Therefore, we get the following result.

Theorem. *Let H be a positive measure. For Ω known, if $p\Omega$ has limiting spectrum H, and $\lim r/p = \kappa$, then the asymptotic misclassification rate $MC^* = \lim_{p\ to\infty} MC$ is a function only of κ, is increasing in κ, and is bounded between*

$$MC^*(0) = \Pr[N(0,1) < M_{K-1}] < 1$$

and

$$MC^*(\infty) = \Pr[N(\mu^*(\infty), 1) < M_{K-1}] > 0,$$

where

$$\mu^*(\infty) = \int \nu^{-1} H(d\nu) = \lim_{p\to\infty} MI.$$

To summarize, the misclassification rate with a fixed number of repeats r converges to a constant, $\mathrm{MC}^*(0) < 1$, while the Bayes misclassification rate is given by $\mathrm{MC}^*(\infty) > 0$. To get a asymptotically constant misclassification rate between $\mathrm{MC}^*(0)$ and $\mathrm{MC}^*(\infty)$, one needs a number of repeats r which is of order $O(p)$.

In the case that one uses $r = \kappa p$ repeats, the asymptotic misclassification rate as a function of K matches the *Bayes* rate of a classification problem with covariance $\tilde{\Omega}$, where the eigenvalues of $\tilde{\Omega}$ are given by $\omega_i + \frac{\omega_i^2}{p\kappa}$; i.e. the eigenvalues of $\tilde{\Omega}$ are inflated compared to Ω. This means that if we were to naively interpret the obtained misclassification rates as Bayes misclassification rates, the estimated mutual information would the true mutual information of $\tilde{\Omega}$, given by

$$\hat{\mathrm{MI}} = \int 1/(\eta^2/\kappa + \eta) H(d\eta).$$

The asymptotic difference between the true mutual information and the naive

estimate is

$$\begin{aligned}\mathrm{MI} - \hat{\mathrm{MI}} &= \int \frac{1}{\eta} - \frac{1}{\eta^2/\kappa + \eta} H(d\eta) \\ &= \int \frac{\eta^2/\kappa}{\eta^2 + \eta^3/\kappa} H(d\eta) \\ &= \int \kappa^{-1} \frac{1}{1 + \eta/\kappa} H(d\eta) \\ &\leq \int \kappa^{-1} H(d\eta) = \kappa^{-1}.\end{aligned}$$

Hence we have a lower bound for the error of MI estimation.

Theorem. *For Ω known, if $p\Omega$ has limiting spectrum H, and $\lim r/p = \kappa$, then the asympotic minimax risk of estimating MI based on the misclassification curve is upper bounded by*

$$\lim_{p\to\infty} \inf_{\hat{\mathrm{MI}}} \mathbf{E}[|\mathrm{MI} - \hat{\mathrm{MI}}|] \leq \frac{1}{\kappa}.$$

2 Appendix

2.1 Gaussian min probs

Define

$$f_{ng}(\mu, \sigma^2, K) = \Pr[\sigma Z_* + \mu < \max_{i=1}^{K} Z_i]$$

for $Z_*, Z_1, \ldots, Z_K$ i.i.d normal.

Suppose

$$\begin{bmatrix} y_* \\ y_1 \\ \vdots \\ y_{K-1} \end{bmatrix} \sim N\left(\begin{bmatrix} a \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \begin{bmatrix} b & c & \ldots & c \\ c & d & \ldots & e \\ \ldots & \ldots & \ddots & \vdots \\ c & e & \ldots & d \end{bmatrix} \right).$$

where $d > e > \frac{c^2}{b}$.

Then

$$\Pr[y_* < \min_{i=1}^{K-1} y_i] = 1 - f_{ng}\left(-\frac{a}{\sqrt{d-e}}, \frac{b+e-2c}{d-e}, K-1 \right).$$