

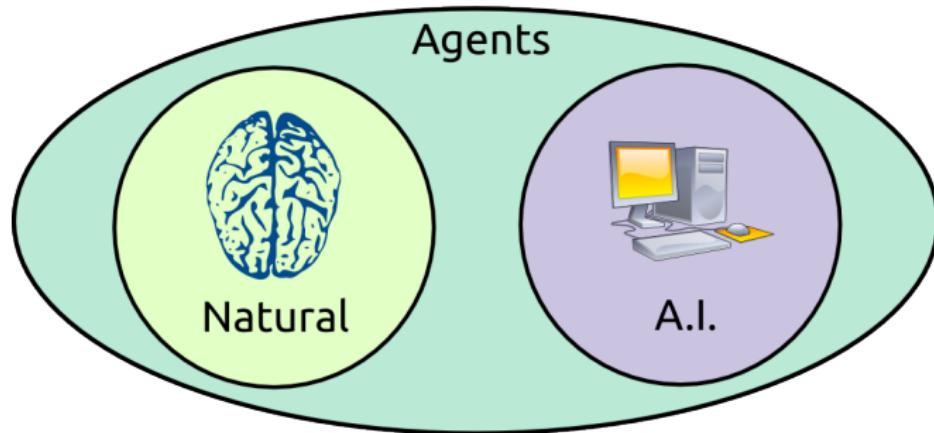
Supervised Evaluation of Representations

Charles Zheng

Stanford University

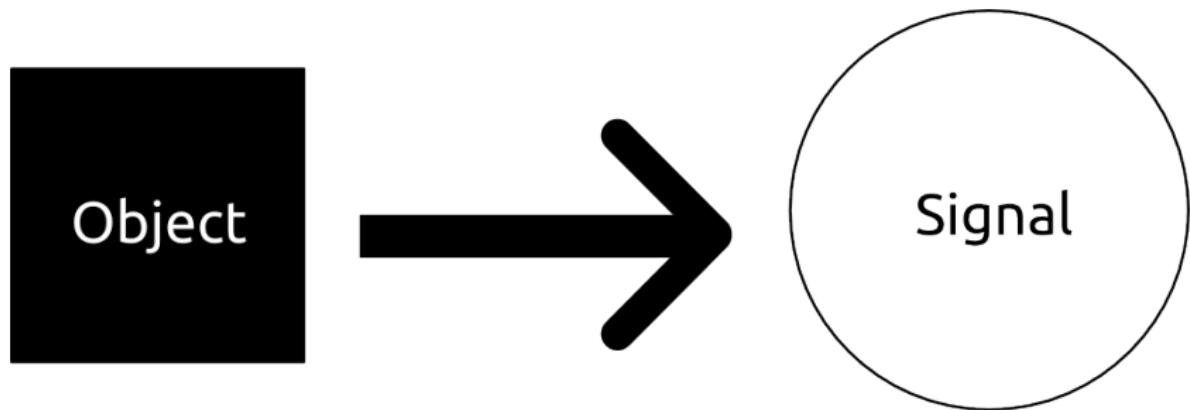
April 26, 2017

Overview



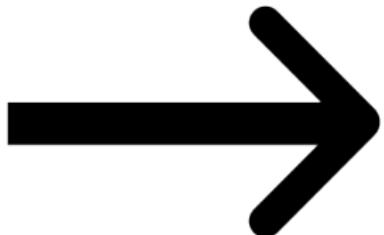
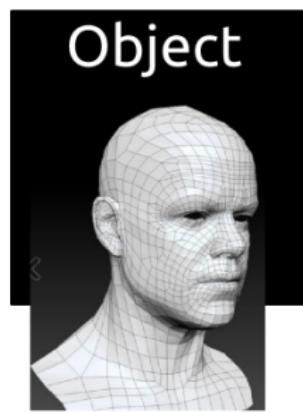
Human brains and machine learning algorithms tackle similar types of problems.

Perception

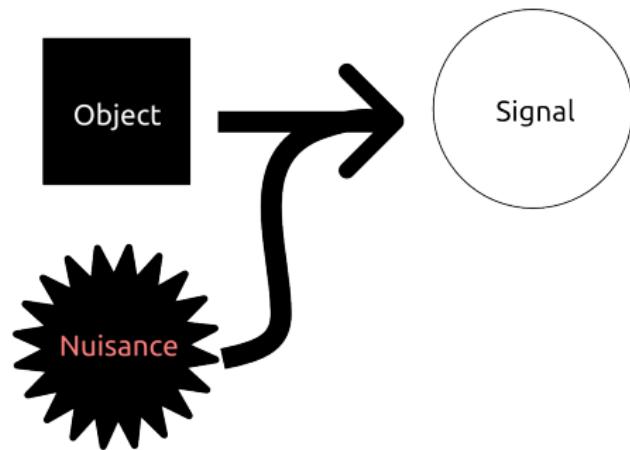


Perception: the problem of inferring *objects* in the environment given observed *signals*.

Example: face recognition

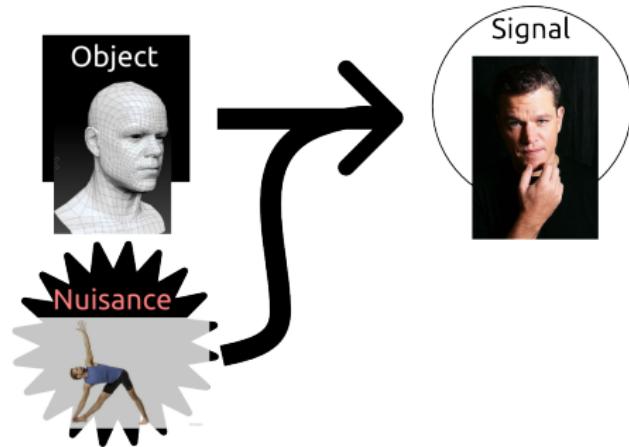


Perception



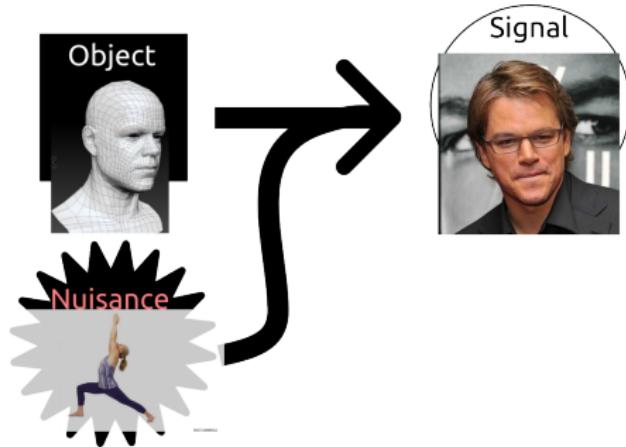
The problem is complicated because there exist some *nuisance parameters*, so the mapping from object to signal is not one-to-one.

Example: face recognition



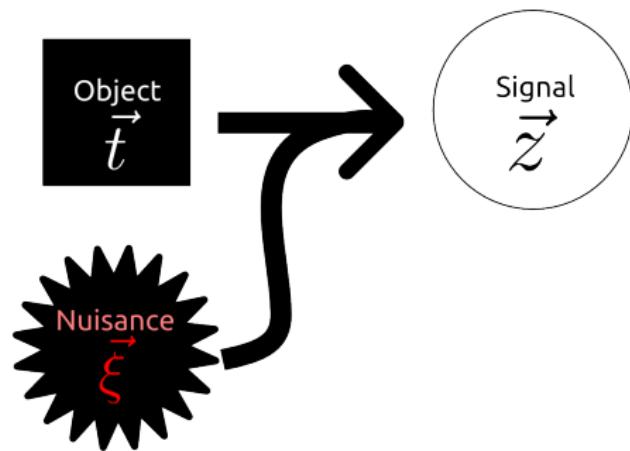
In face recognition, the *pose* (including hairstyle) and *lighting* are nuisance parameters.

Example: face recognition



The same object can map to multiple signals.

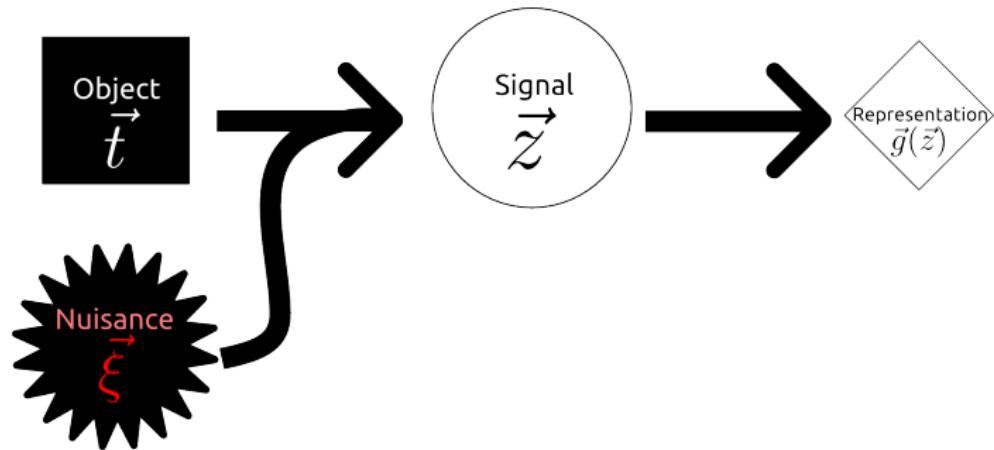
Perception



Assume there exists a function ψ that maps objects and nuisance parameters to signals:

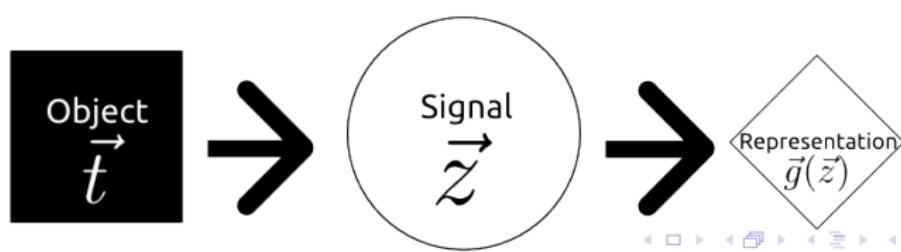
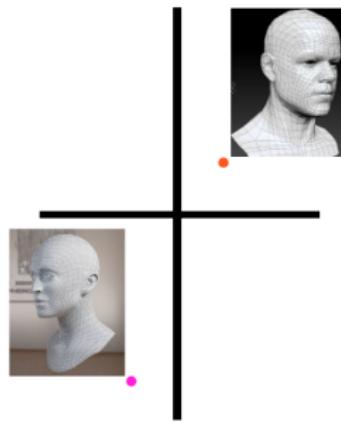
$$\vec{z} = \psi(\vec{t}, \vec{\xi}).$$

What is a representation?

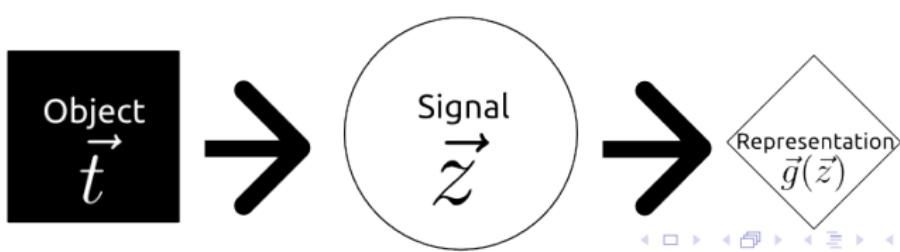
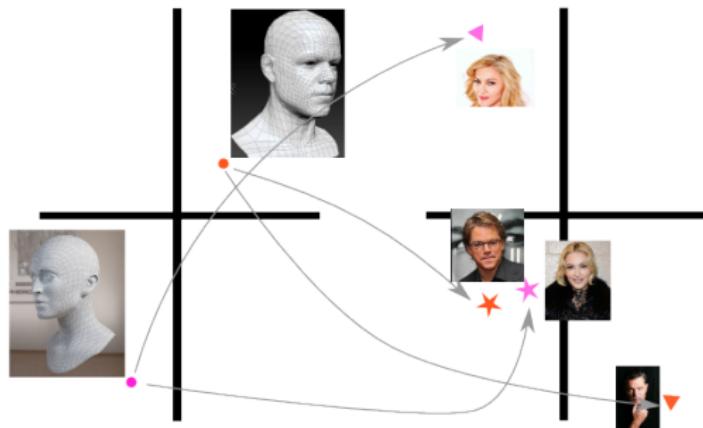


A dimensionality-reducing mapping \vec{g} of the signal.

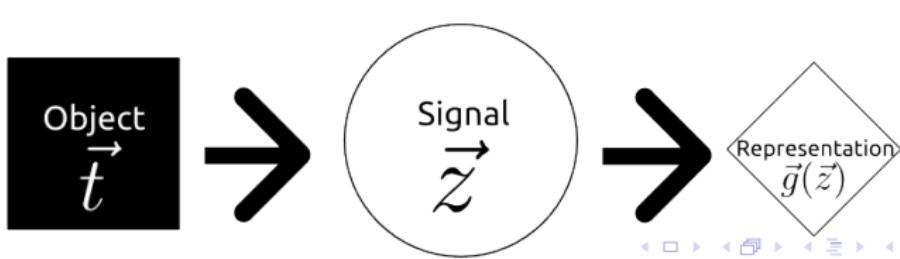
A good representation...



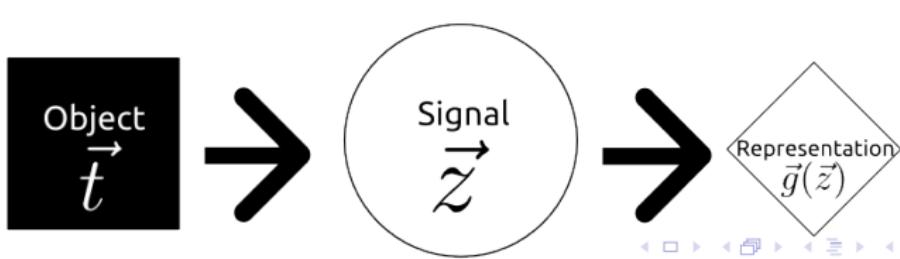
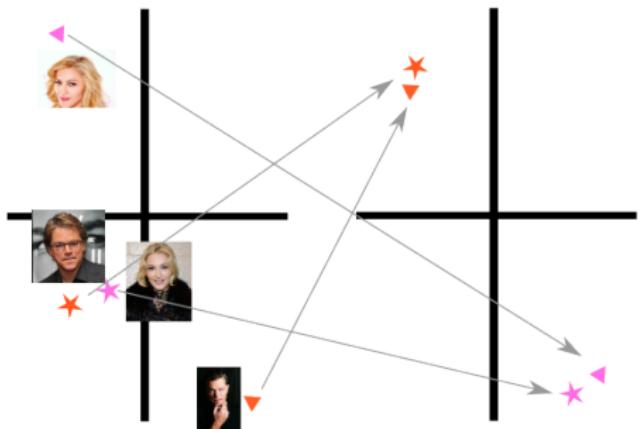
A good representation...



A good representation...



...captures the object space geometry



Why do we care?

- 1) Neuroscience. The brain is hypothesized to use representations for cognitive purposes

Why do we care?

- 1) Neuroscience. The brain is hypothesized to use representations for cognitive purposes
- 2) Machine learning. Representations turn out to be useful for many Machine Learning tasks!

How can we tell if a representation is good?

- Method 1: *Ground truth*. If we happen to know the object parameters \vec{t} (e.g. we simulated the data).

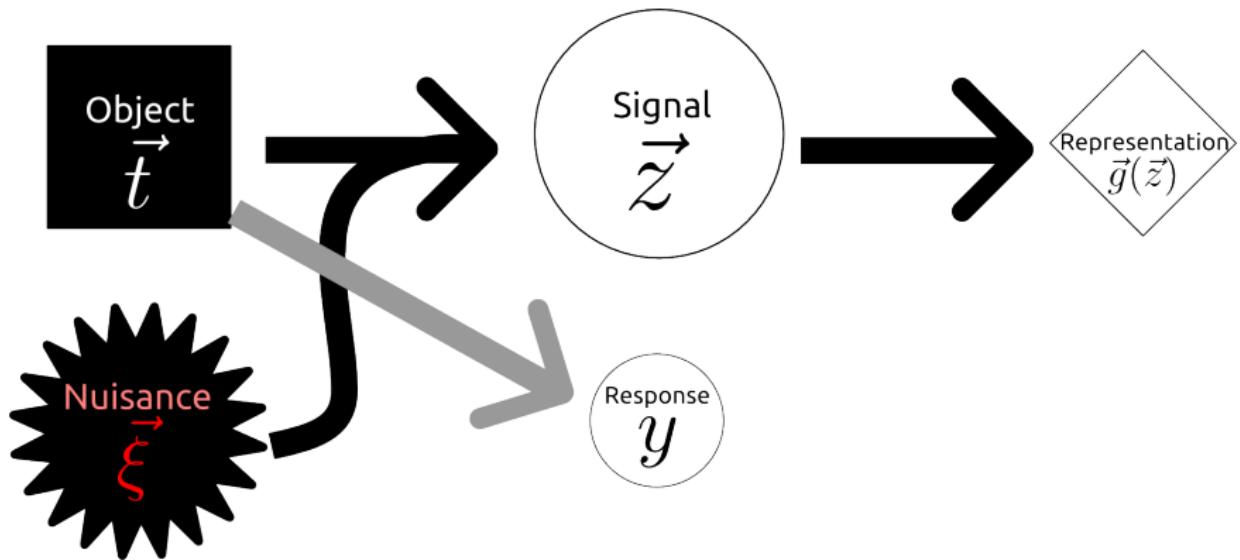
How can we tell if a representation is good?

- Method 1: *Ground truth*. If we happen to know the object parameters \vec{t} (e.g. we simulated the data).
- Method 2: *End result*. By the performance of the representation on a machine learning task.

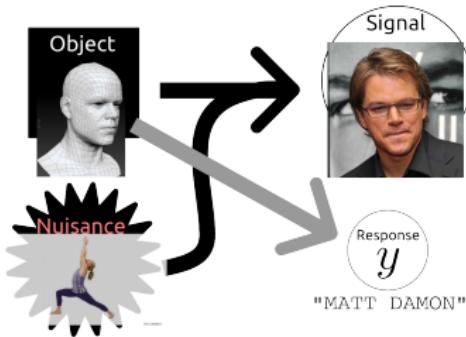
How can we tell if a representation is good?

- Method 1: *Ground truth*. If we happen to know the object parameters \vec{t} (e.g. we simulated the data).
- Method 2: *End result*. By the performance of the representation on a machine learning task.
- Method 3: *Supervised*. If we have a *response variable* Y which can be used to infer distances in \vec{t} .

Supervised evaluation of representations

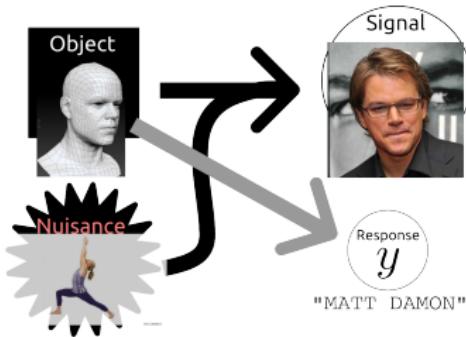


Example: face recognition



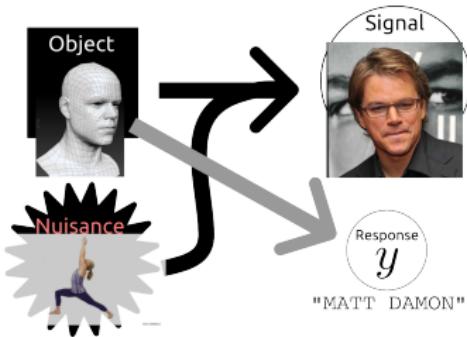
- The ID of the individual is an appropriate *response* variable...

Example: face recognition



- The ID of the individual is an appropriate *response* variable...
- ...because two photos labeled with the same ID must belong to the same object \vec{t}

Example: face recognition



- The ID of the individual is an appropriate *response* variable...
- ...because two photos labeled with the same ID must belong to the same object \vec{t}
- That is, for $d(y, y')$ being the zero-one distance,

$$d(y, y') = 0 \Leftrightarrow d(\vec{t}, \vec{t}') = 0.$$

Outline

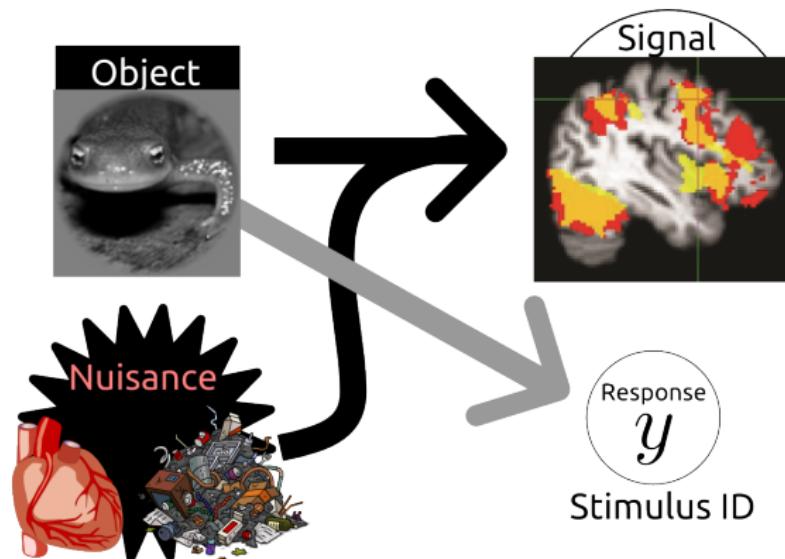
- *Previous work.* Identification accuracy, a method for evaluating representations
- *Contribution 1.* Extrapolation of identification accuracy.
- *Contribution 2.* Link between identification accuracy and mutual information.

Section 2

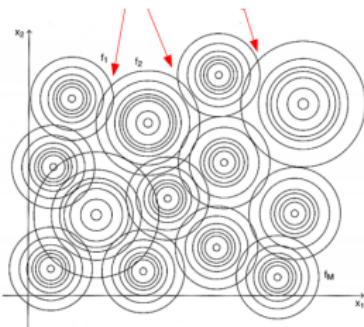
Identification

Identifying natural images from fMRI data

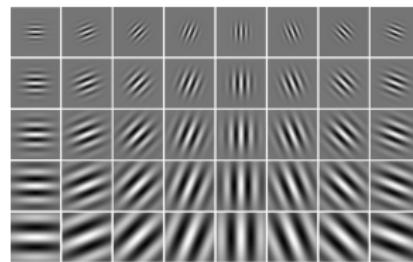
Kay, Naselaris, Prenger and Gallant (2008), *Nature*.



Comparing two different natural image bases



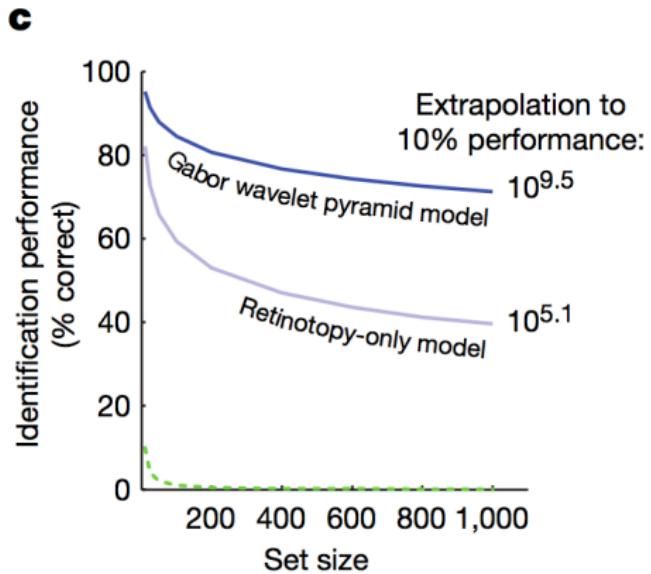
Retinotopic (Gaussian) basis



Gabor filter basis

Identifying natural images from fMRI data

Kay et al. used *identification accuracy* as a metric for the quality of the representation



Multiple-response regression

- Pairs $(x_i, y_i)_{i=1}^n$, where X is p -dimensional and Y is q -dimensional.
- Data matrices $\mathbf{X}_{n \times p}$, $\mathbf{Y}_{n \times q}$.
- For each column of Y , fit sparse model $Y^{(i)} \approx X^T \beta^{(i)} + \epsilon$, e.g. by using elastic net (Zou 2008),

$$\hat{\beta}^{(i)} = \operatorname{argmin}_{\beta} \|\mathbf{X}^T \beta^{(i)} - Y^{(i)}\|^2 + \lambda_2 \|\beta^{(i)}\|_2^2 + \lambda_1 \|\beta^{(i)}\|_1$$

Regression vs Identification accuracy

- Independent test set $(x_i^*, y_i^*)_{i=1}^k$.
- Use model to predict $\hat{y}_i^* = (x_i^*)^T \hat{B}$ for $i = 1, \dots, k$.

Two ways to evaluate the predictive accuracy of the regression model:

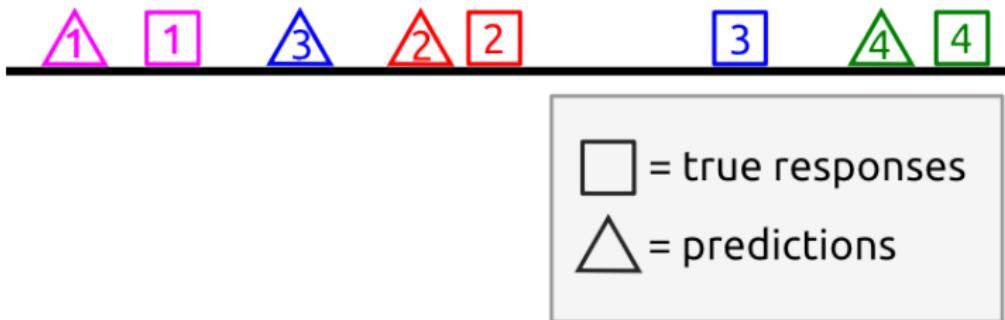
- Regression (mean squared-error) loss:

$$\text{MSE} = \frac{1}{k} \sum_{i=1}^k \|y_i^* - \hat{y}_i^*\|^2.$$

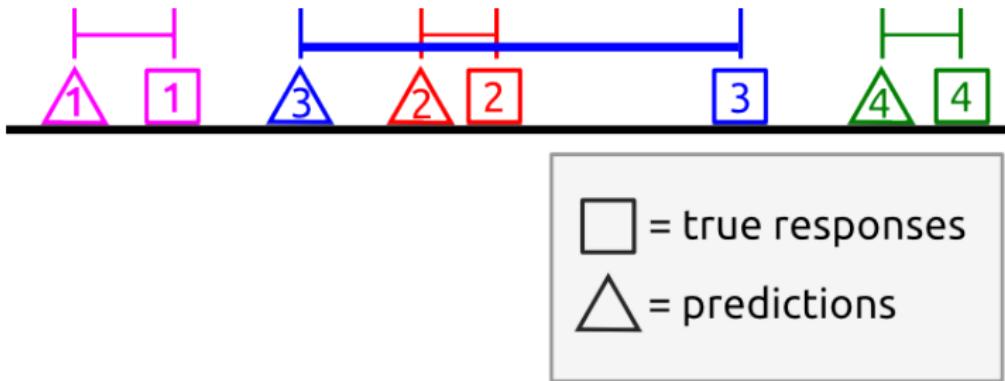
- Identification accuracy (Kay 2008):

$$\text{IdAcc}_k = \frac{1}{k} \sum_{i=1}^k I\{\hat{y}_i^* \text{ is nearest neighbor of } y_i^*\}.$$

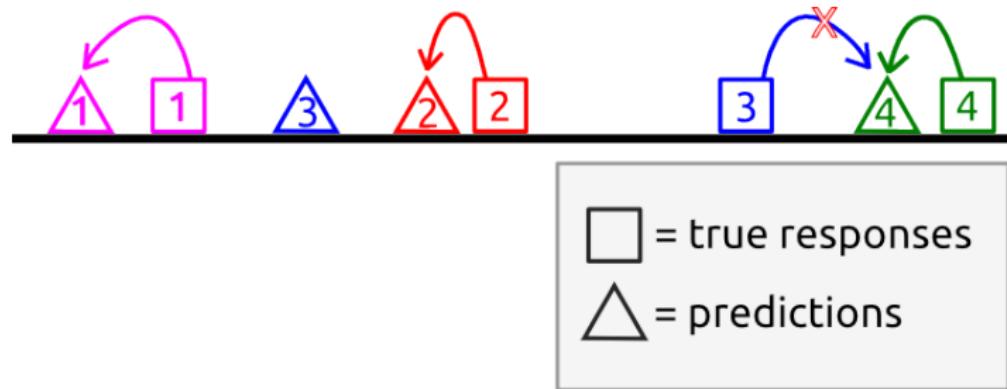
Regression vs Identification accuracy



Mean-squared error



Identification accuracy



Identification accuracy for comparing representations

- Suppose you have two different representation models:
 - \vec{g}_1 , the retinotopic model
 - \vec{g}_2 , Gabor filters

Identification accuracy for comparing representations

- Suppose you have two different representation models:
 - \vec{g}_1 , the retinotopic model
 - \vec{g}_2 , Gabor filters
- For each model, train a linear model

$$\vec{Y} = B^T \vec{g}(\vec{Z}) + \epsilon$$

Identification accuracy for comparing representations

- Suppose you have two different representation models:
 - \vec{g}_1 , the retinotopic model
 - \vec{g}_2 , Gabor filters
- For each model, train a linear model

$$\underbrace{\vec{Y}}_{\text{brain response}} = \underbrace{B^T}_{\text{coefficient matrix}} \underbrace{\vec{g}}_{\text{representation}} (\underbrace{\vec{Z}}_{\text{pixels}}) + \epsilon$$

Identification accuracy for comparing representations

- Suppose you have two different representation models:
 - \vec{g}_1 , the retinotopic model
 - \vec{g}_2 , Gabor filters
- For each model, train a linear model

$$\underbrace{\vec{Y}}_{\text{brain response}} = \underbrace{B^T}_{\text{coefficient matrix}} \underbrace{\vec{g}}_{\text{representation}} (\underbrace{\vec{Z}}_{\text{pixels}}) + \epsilon$$

- Evaluate the identification accuracy on k test images.

Identification accuracy for comparing representations

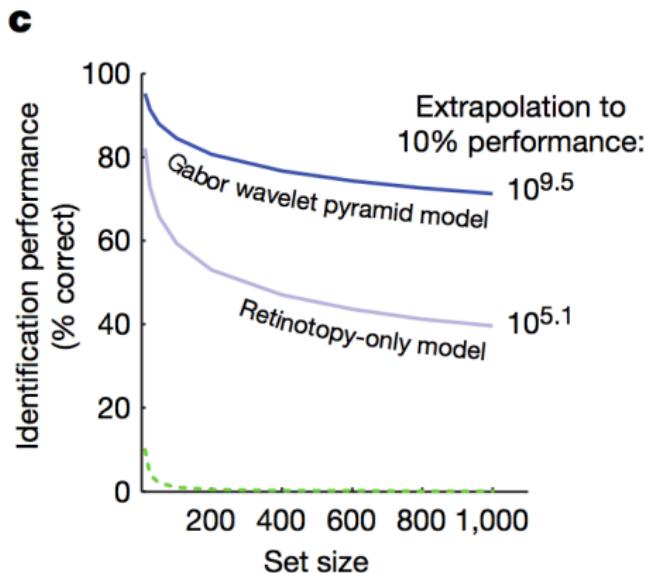
- Suppose you have two different representation models:
 - \vec{g}_1 , the retinotopic model
 - \vec{g}_2 , Gabor filters
- For each model, train a linear model

$$\underbrace{\vec{Y}}_{\text{brain response}} = \underbrace{B^T}_{\text{coefficient matrix}} \underbrace{\vec{g}}_{\text{representation}} (\underbrace{\vec{Z}}_{\text{pixels}}) + \epsilon$$

- Evaluate the identification accuracy on k test images.
- You can use a test set which is larger than k , and average the identification accuracy over k -sized subsamples

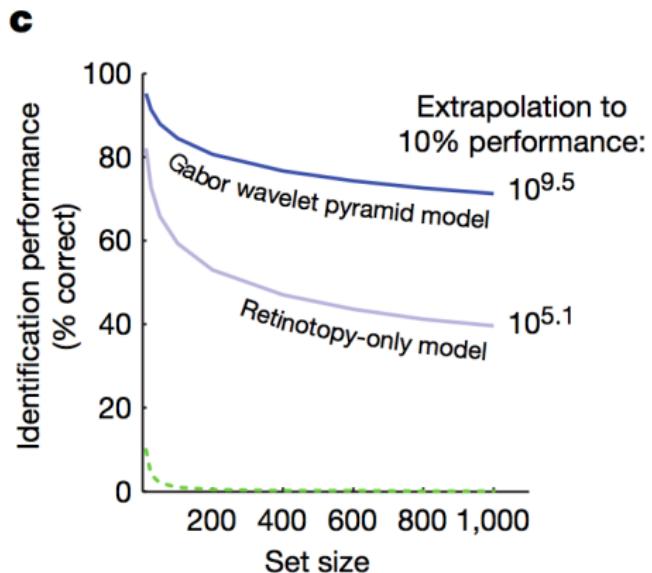
Identifying natural images from fMRI data

Gabor filters yields consistently higher accuracy than retinotopic model



Identifying natural images from fMRI data

Gabor filters yields consistently higher accuracy than retinotopic model



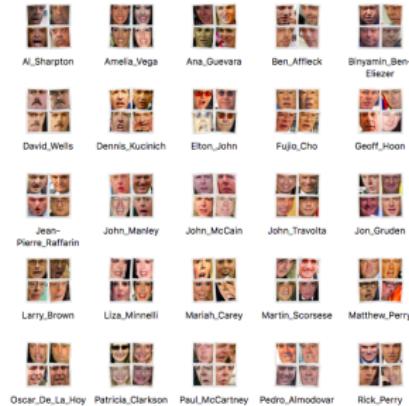
Q: Is this always the case? or could you also have intersecting curves?

Section 3

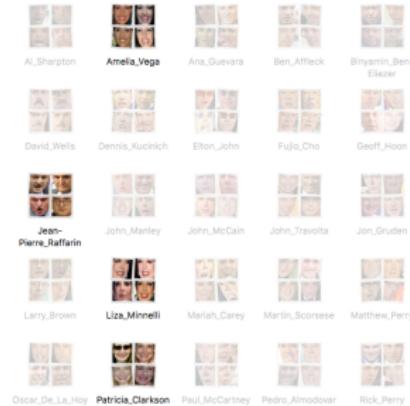
Extrapolation

Randomized multi-class classification

1. Population of categories $\pi(y)$

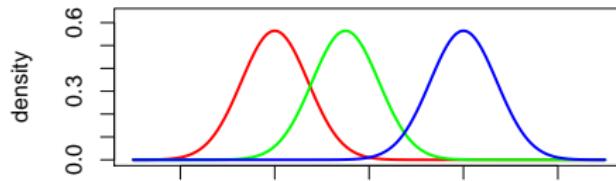


2. Subsample k labels, y_1, \dots, y_k



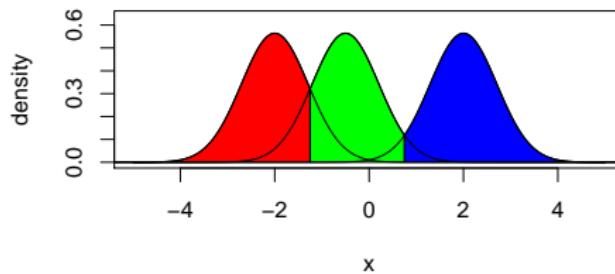
Identification is a special case of a *randomized classification* task.

Toy example



- Suppose $k = 3$, and we draw Y_1, Y_2, Y_3 .
- The *Bayes rule* is the optimal classifier and depends on knowing the true densities:
$$\hat{y}(x) = \operatorname{argmax}_{y_i} p(x|y_i)$$
- The *Bayes Risk*, which is the misclassification rate of the optimal classifier.

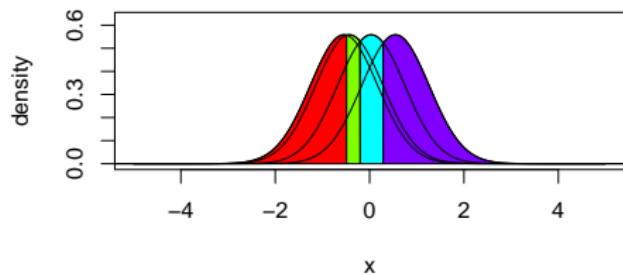
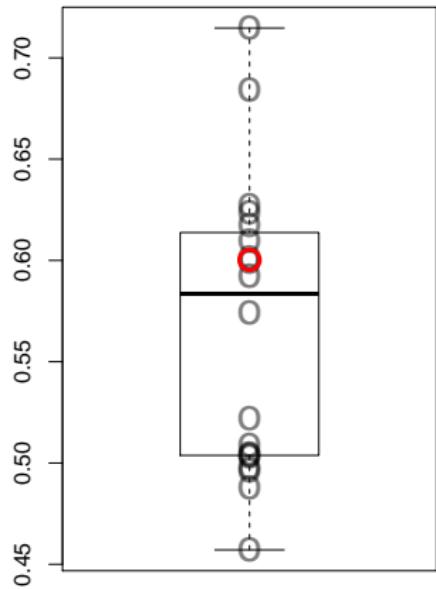
Toy example



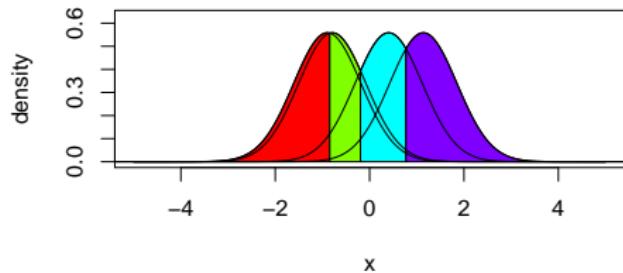
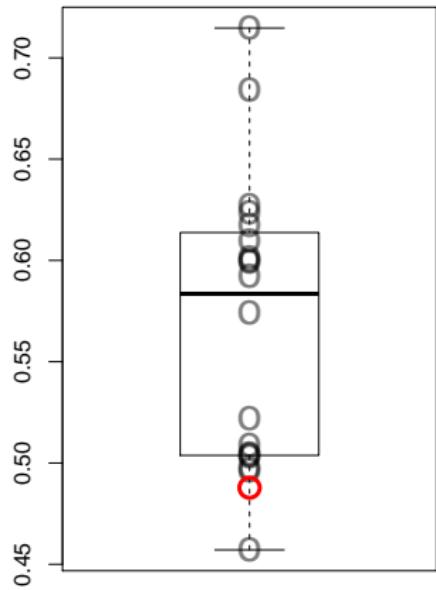
- The *Bayes Risk* is the expected test error of the Bayes rule,

$$\frac{1}{k} \sum_{i=1}^k \Pr[\hat{y}(x) \neq Y | Y = y_i]$$

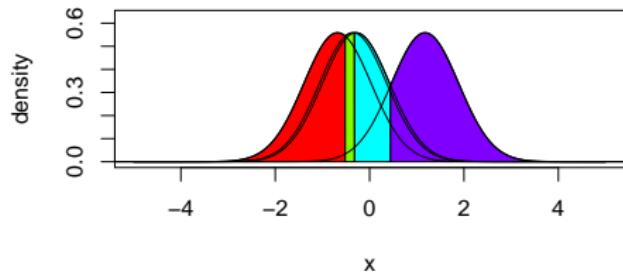
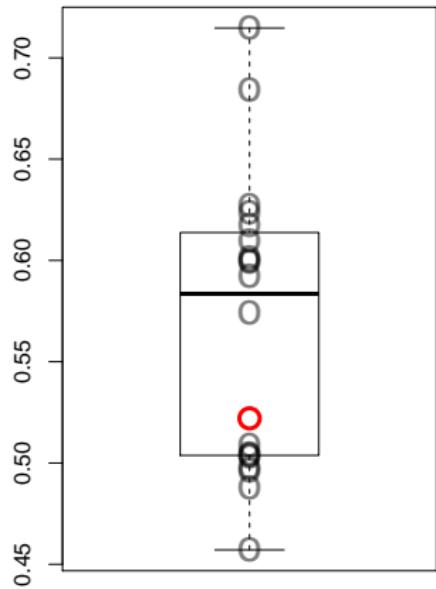
Toy example



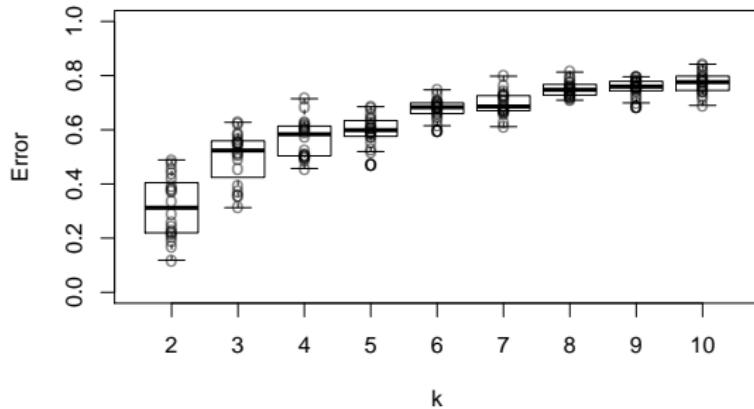
Toy example



Toy example

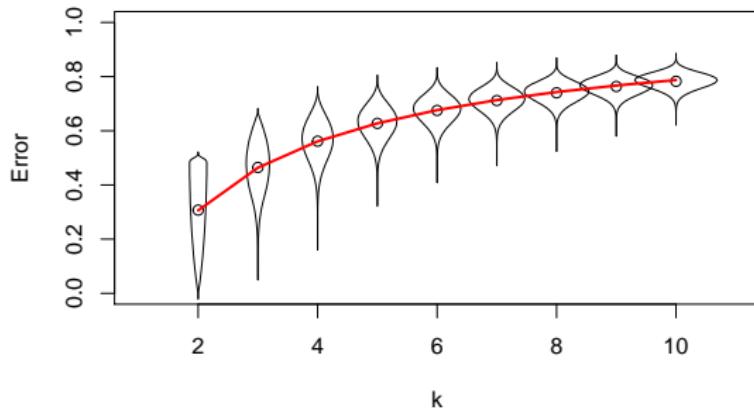


Toy example



Toy example

$\rho = 0.7$



Theoretical Result

Theorem. (Z., Achanta, Benjamini.) Suppose π , $\{F_y\}_{y \in \mathcal{Y}}$ and marginal classifier \mathcal{F} satisfy (*some regularity condition*). Then, there exists some function $\bar{D}(u)$ on $[0, 1] \rightarrow [0, 1]$ such that the k -class average risk is given by

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$

Theoretical Result

Theorem. (Z., Achanta, Benjamini.) Suppose π , $\{F_y\}_{y \in \mathcal{Y}}$ and marginal classifier \mathcal{F} satisfy (*some regularity condition*). Then, there exists some function $\bar{D}(u)$ on $[0, 1] \rightarrow [0, 1]$ such that the k -class average risk is given by

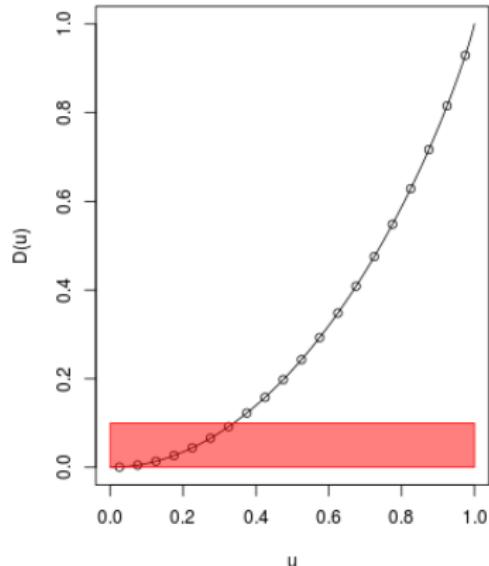
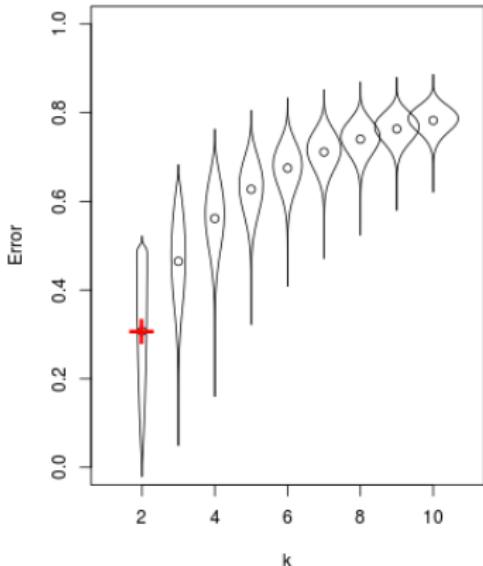
$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$

Remark. This result also applies to the *Bayes identification risk*.

Computing average risk

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$

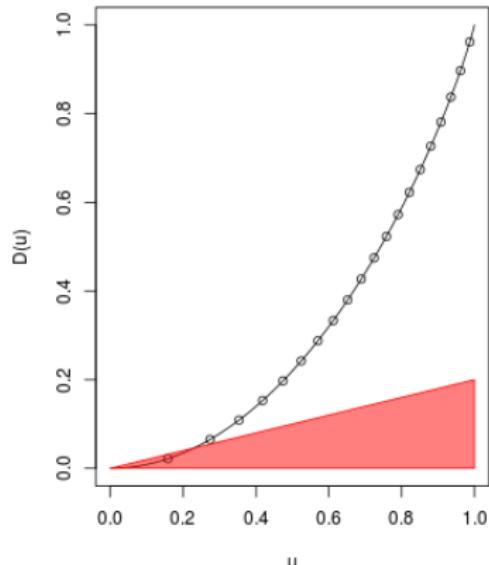
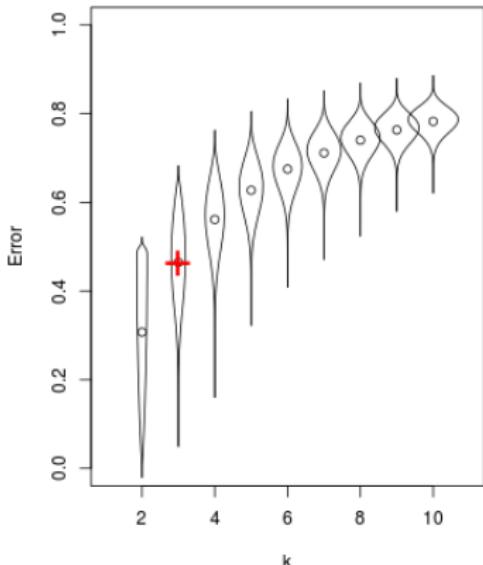
$(k = 2)$



Computing average risk

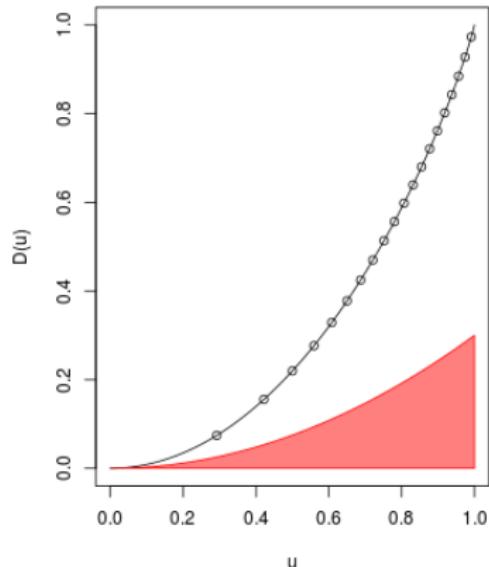
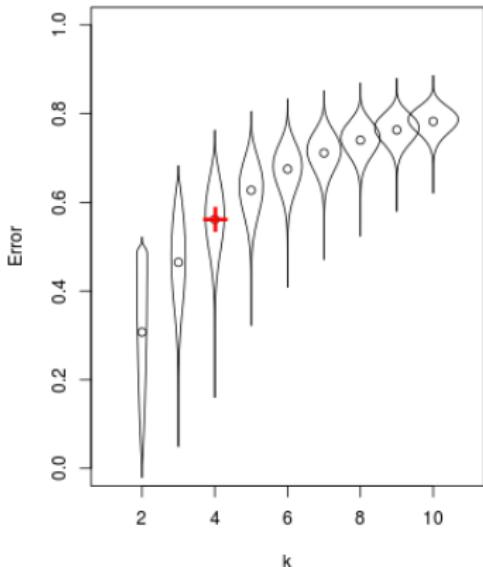
$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$

$(k = 3)$



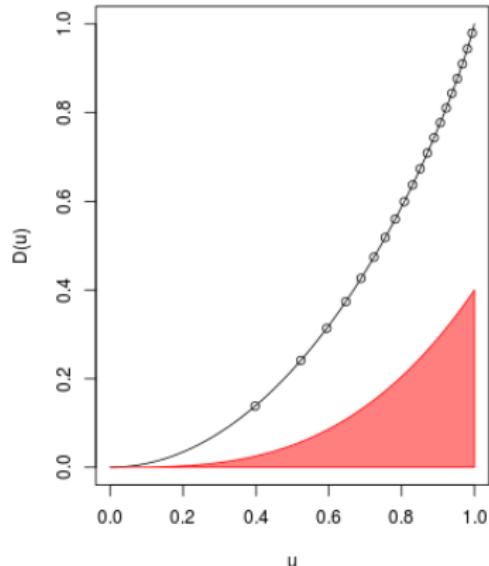
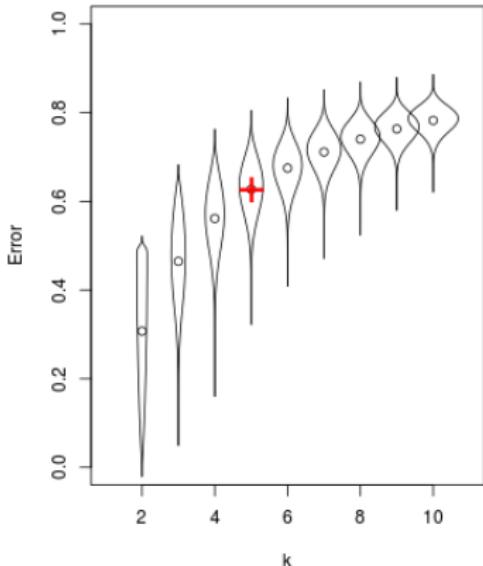
Computing average risk

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$



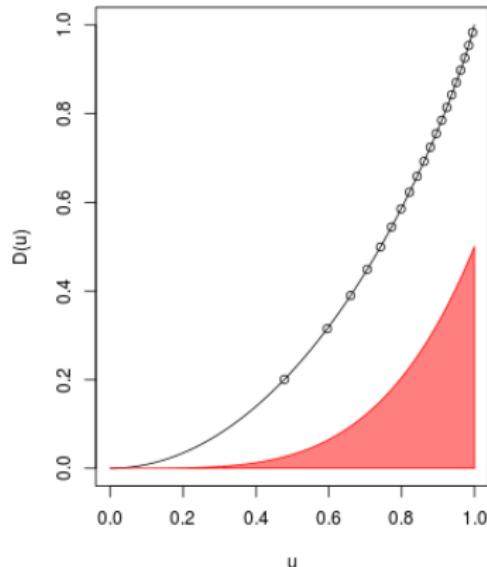
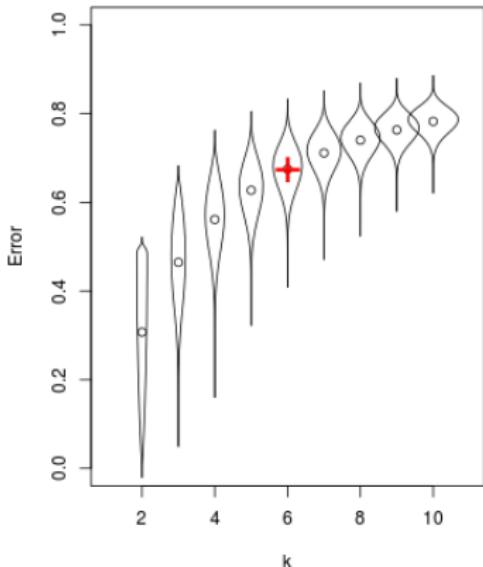
Computing average risk

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$



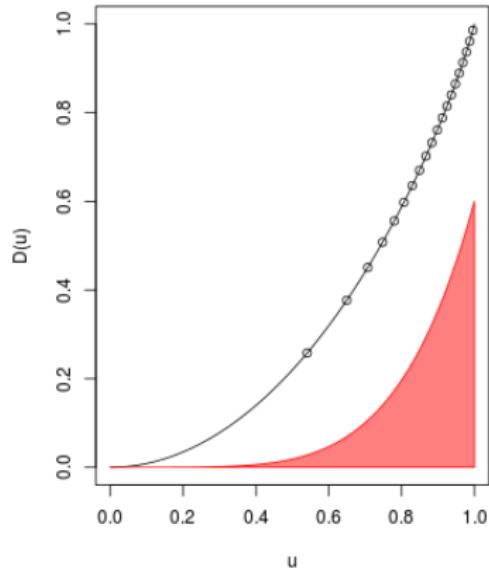
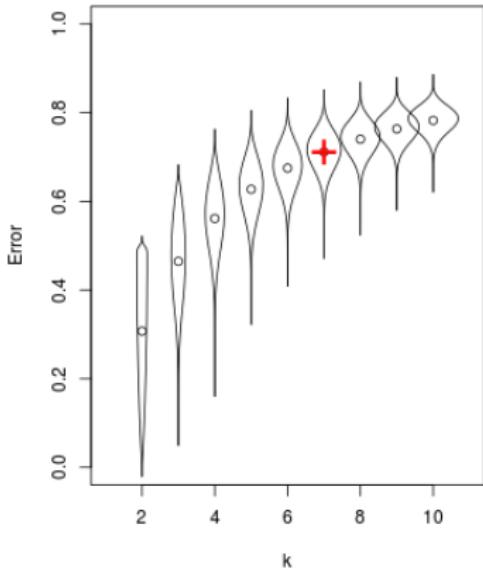
Computing average risk

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$



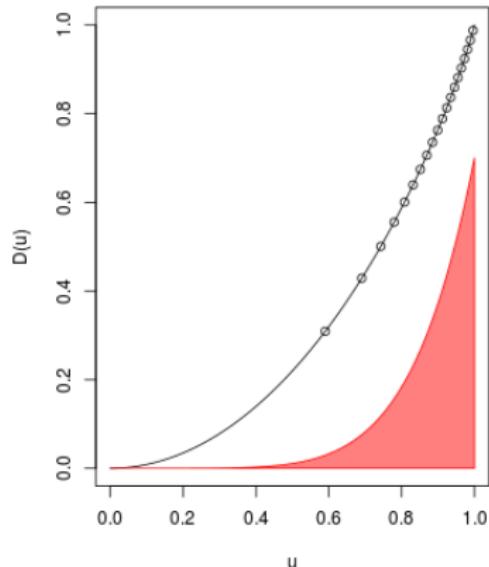
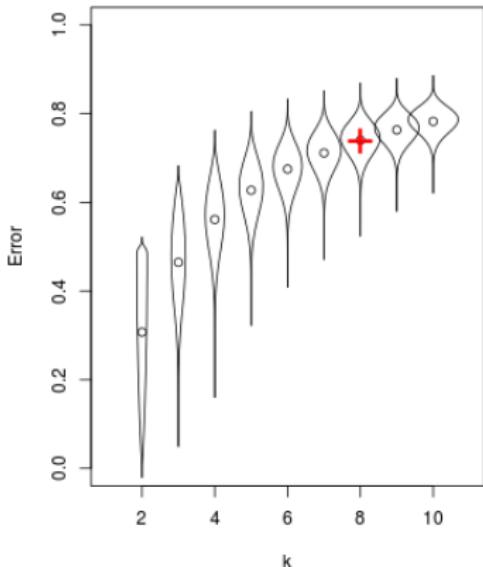
Computing average risk

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$



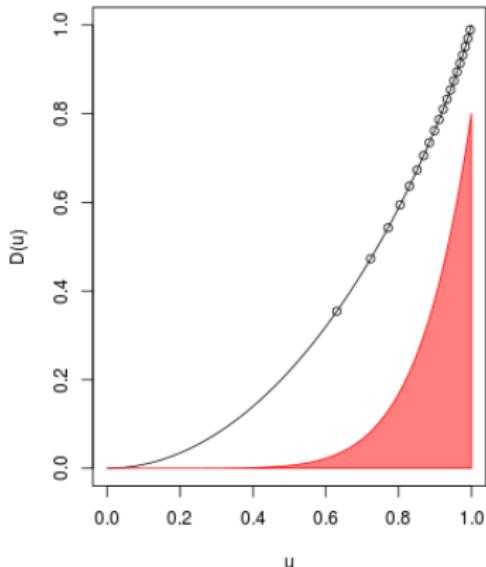
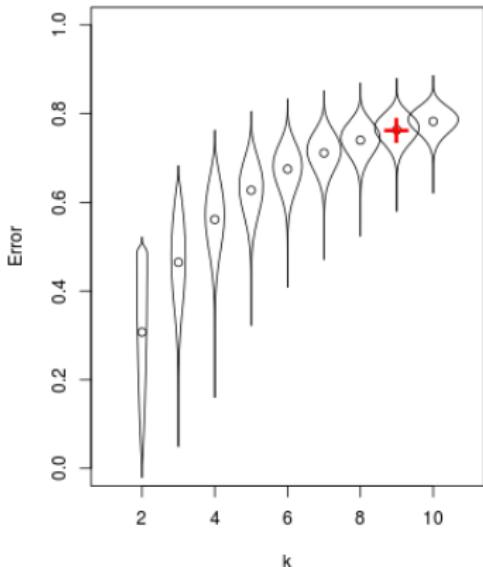
Computing average risk

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$



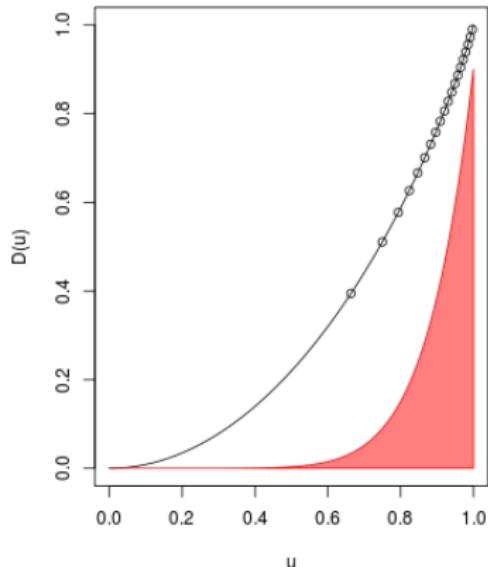
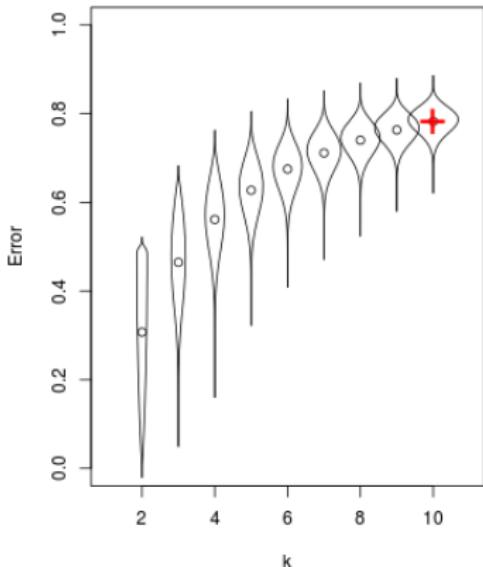
Computing average risk

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$



Computing average risk

$$\text{AvRisk}_k = (k - 1) \int \bar{D}(u) u^{k-2} du.$$



Implication: estimate $\bar{D}(u)$ to predict risk

- Theoretical result links k -class average risk to $\bar{D}(u)$ function
- In real data, we do not know $\bar{D}(u)$ since it depends on the unknown joint distribution
- However, given a model, we can estimate $\bar{D}(u)$

So... can accuracy curves intersect?

- In general, the answer is yes.

So... can accuracy curves intersect?

- In general, the answer is yes.
- However, we will see in the next section that under *high-dimension* assumptions, the Bayes accuracy curves do *not* intersect.

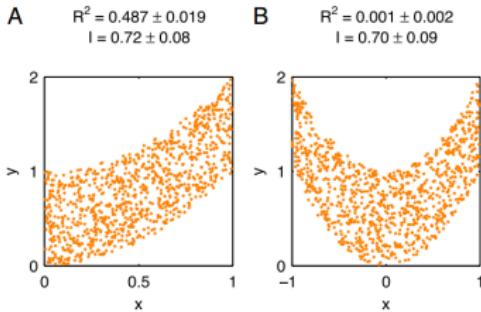
So... can accuracy curves intersect?

- In general, the answer is yes.
- However, we will see in the next section that under *high-dimension* assumptions, the Bayes accuracy curves do *not* intersect.
- Therefore, a single parameter, the *mutual information*, suffices to summarize the entire curve.

Section 4

Mutual Information

Mutual information $I(X; Y)$



Introduced in Shannon's 1948 paper, "A mathematical theory of communication"

$$I(X; Y) = \int \log \left(\frac{p(x, y)}{p(x)p(y)} \right) p(x, y) dx dy$$

Image credit Kinney et al. 2014.

Result 1. Lower bound for mutual information

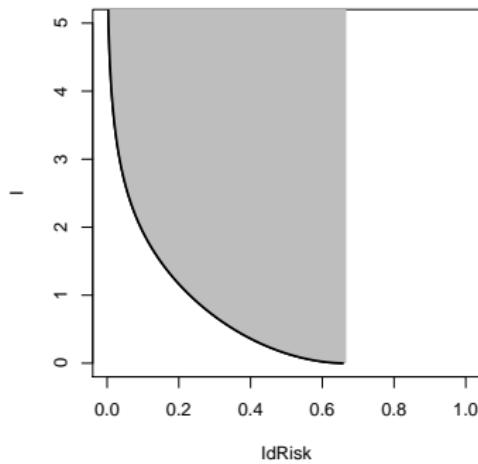
- Define the identification risk as the expected identification loss

$$\text{IdRisk}_k = \mathbf{E}[\text{IdLoss}_k]$$

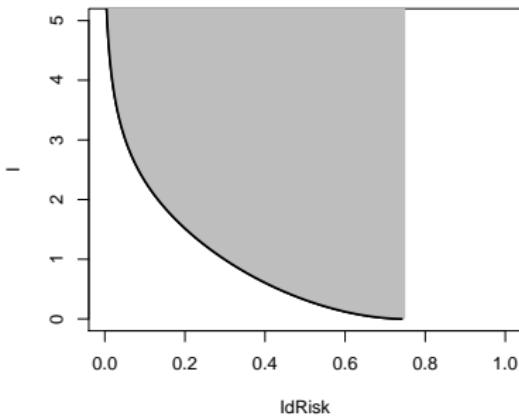
- Theorem.** (Z., Benjamini 2017) There exists a function h_k such that

$$I(\vec{g}(\vec{Z}); \vec{Y}) \geq h_k(\text{IdRisk}_k).$$

h_3



h_4



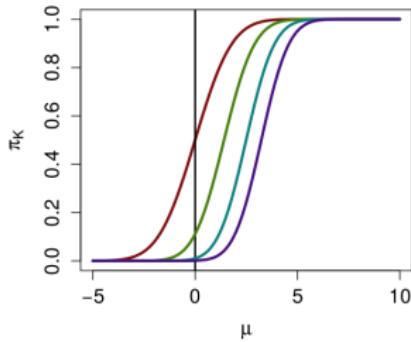
Result 2. Limiting behavior of accuracy curves

(Z., Benjamini 2016) Define ABA_k as the Bayes identification accuracy (or average Bayes classification accuracy). Then under a particular high-dimensional limit,

$$\text{ABA}_k \approx \pi_k(\sqrt{2I(X; Y)}) \quad (1)$$

The function π_k is given by

$$\pi_k(c) = \int_{\mathbb{R}} \phi(z - c) \Phi(z)^{k-1} dz.$$

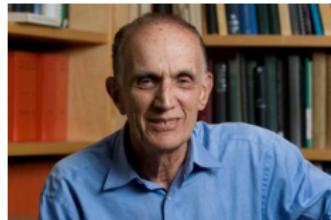


Legend: $K = \{ \boxed{2}, \boxed{9}, \boxed{99}, \boxed{999} \}$

Section 5

Acknowledgements







Section 6

The end