How many neurons does it take to classify a lightbulb?

Charles Zheng

Stanford University

December 14, 2015

(Joint work with Yuval Benjamini)

Overview

Background and motivation

- Review of information theory
- Study of face recognition

Problems

- Estimating mutual information between stimulus and response.
- Can we use machine learning methods to estimate MI?

Methods

- Setup
- Gaussian example
- Using Fano's inequality
- Using low-SNR universality

Results

Information theory

The complexity of modern communications system is made possible by Shannon's theory of information.

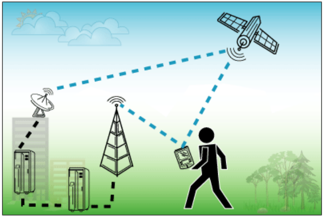

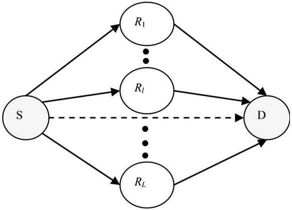

A information-processing network can be analyzed in terms of interactions between its components (which are viewed as random variables. Image credit CartouCHe, Aziz et al. 2011

Entropy and mutual information

X and Y have joint density $p(x, y)$ with respect to μ.

Quantity	Definition	Linear analogue
Entropy	$H(X) = -\int (\log p(x)) p(x) \mu_X(dx)$	$\text{Var}(X)$
Conditional entropy	$H(X\|Y) = \mathbf{E}[H(X\|Y)]$	$\mathbf{E}[\text{Var}(X\|Y)]$
Mutual information	$I(X; Y) = H(X) - H(X\|Y)$	$\text{Cor}(X, Y)$

The above definition includes both *differential* entropy and *discrete* entropy. Information theorists tend to use log base 2, we will use natural logs in this talk.

Properties of mutual information

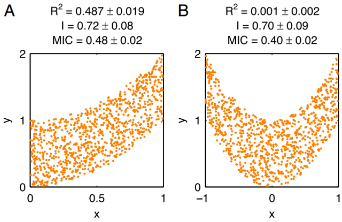

- Nonnegative: $I(X;Y) \geq 0$
- Symmetric: $I(X;Y) = I(Y;X)$
- Bijection-invariant: $I(\phi(X);\psi(Y)) = I(\psi(Y);\phi(X)$.
- Additivity. If $(X_1, Y_1) \perp (X_2, Y_2)$, then

$$I((X_1, X_2); (Y_1, Y2)) = I(X_1; Y_1) + I(X_2; Y_2)$$

- Relation to KL divergence $\mathbb{D}$.

$$\mathbb{D}(p(x,y)||p(x)p(y)) = I(X;Y)$$

Image credit Kinney et al. 2014

Studying the neural code

The brain is the *most complex* information processing system we know!

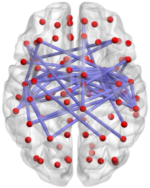

Neural network inferred from data (Hong et al.)

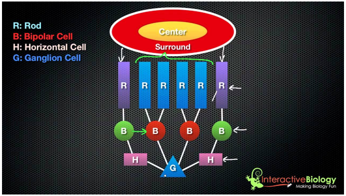

Organization of human retina

How do neurons encode, process, and decode sensory information?

Image credit: Hong et al., Interactive Biology

Studying the neural code: data

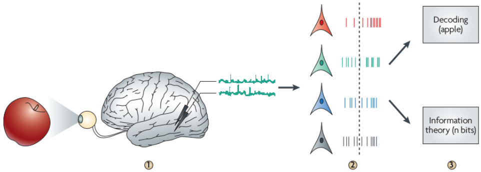

- Let $\mathcal{X}$ define a class of stimuli (faces, objects, sounds.)
- Stimulus $\mathbf{X} = (X_1, \ldots, X_p)$, where X_i are features (e.g. pixels.)
- Present $\mathbf{X}$ to the subject, record the subject's brain activity using EEG, MEG, fMRI, or calcium imaging.
- Recorded response $\mathbf{Y} = (Y_1, \ldots, Y_q)$, where Y_i are single-cell responses, or recorded activities in different brain region.

Image credits: Quiroga et al. (2009)

Experimental setup

- How to make inferences about the population of stimuli in $\mathcal{X}$ using finitely many examples?
- *Randomization.* Select $\mathbf{X}^{(1)}, \ldots, \mathbf{X}^{(K)}$ randomly from some distribution $p(\mathbf{x})$ (e.g. an image database). Record r responses from each stimulus.

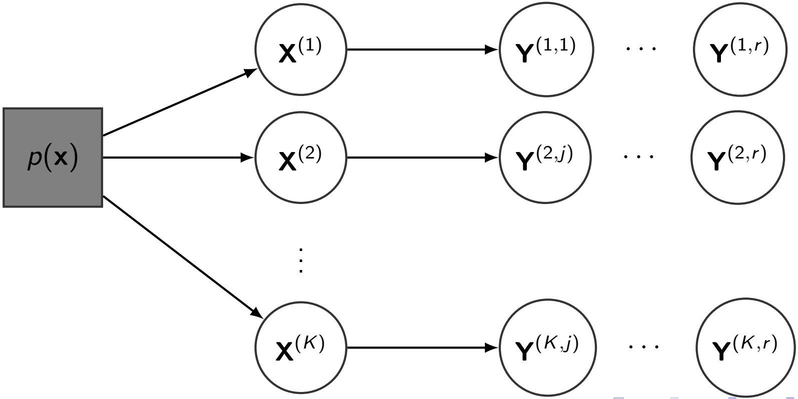

References

- Cover and Thomas. Elements of information theory.
- Muirhead. Aspects of multivariate statistical theory.
- van der Vaart. Asymptotic statistics.